

SCM-516

Team Final Project-Team 407

Team Members

1. Anwesha Banerjee
2. Qinxiin Zhang
3. Alvin Lee
4. Satyanand Eranki
5. Xuzhou Feng

Classification

Dataset



Bank Marketing

Donated on 2/13/2012

The data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution. The classification goal is to predict if the client will subscribe a term deposit (variable y).

Dataset Characteristics

Multivariate

Subject Area

Business

Associated Tasks

Classification

Feature Type

Categorical, Integer

Instances

45211

Features

16

- Moro, S., Rita, P., & Cortez, P. (2014). Bank Marketing [Dataset]. UCI Machine Learning Repository. <https://doi.org/10.24432/C5K306>.

Data Preprocessing

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y
2424	29	blue-collar	single	secondary	no	783	yes	no	unknown	13	may	166	1	-1	0	unknown	no
34231	33	admin.	married	tertiary	no	69	no	yes	cellular	4	may	77	2	-1	0	unknown	no
26140	40	entrepreneur	married	secondary	no	1349	no	no	cellular	20	nov	82	3	-1	0	unknown	no
25200	36	services	single	secondary	no	461	no	no	cellular	18	nov	434	2	-1	0	unknown	no
9167	48	technician	married	unknown	no	460	no	no	unknown	5	jun	99	3	-1	0	unknown	no

Output after Data Processing-Part 1

	age	job	marital	education	default	balance	housing	loan	duration	campaign	y
23787	47	blue-collar	married	primary	no	961	no	no	45	11	no
36427	37	technician	single	tertiary	no	555	yes	no	180	3	no
1547	53	blue-collar	married	secondary	no	359	no	yes	238	8	no
33652	29	blue-collar	married	secondary	no	457	yes	no	72	3	no
27706	30	blue-collar	single	tertiary	no	4127	yes	no	47	4	no

Data Preprocessing contd...

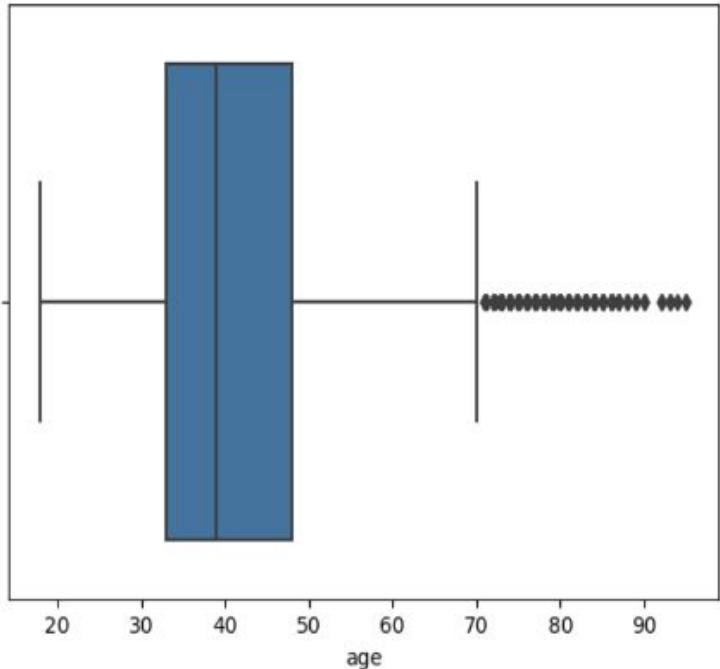
Checking for “Missing Values”

```
age          0  
job          0  
marital      0  
education    0  
default      0  
balance      0  
housing      0  
loan         0  
duration     0  
campaign     0  
y             0  
dtype: int64
```

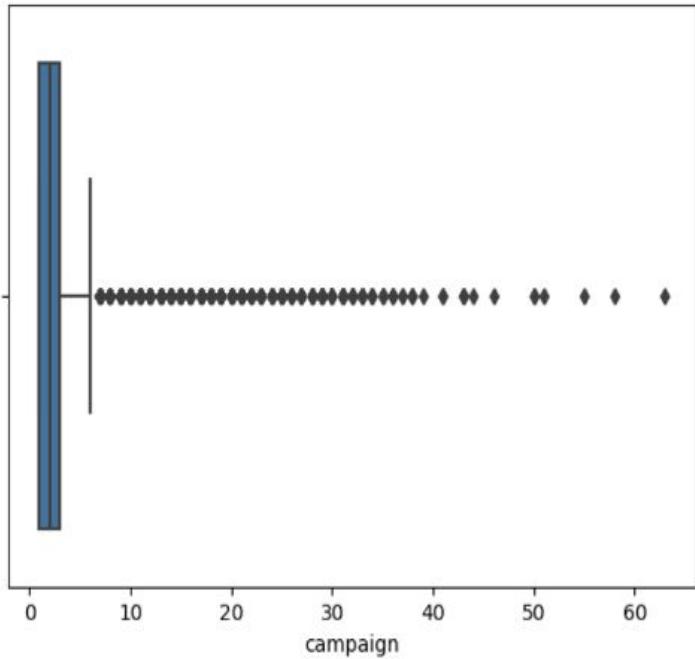
In our data set it is seen that we do not have any “missing values”

Data Preprocessing contd...

Removing “Outliers”



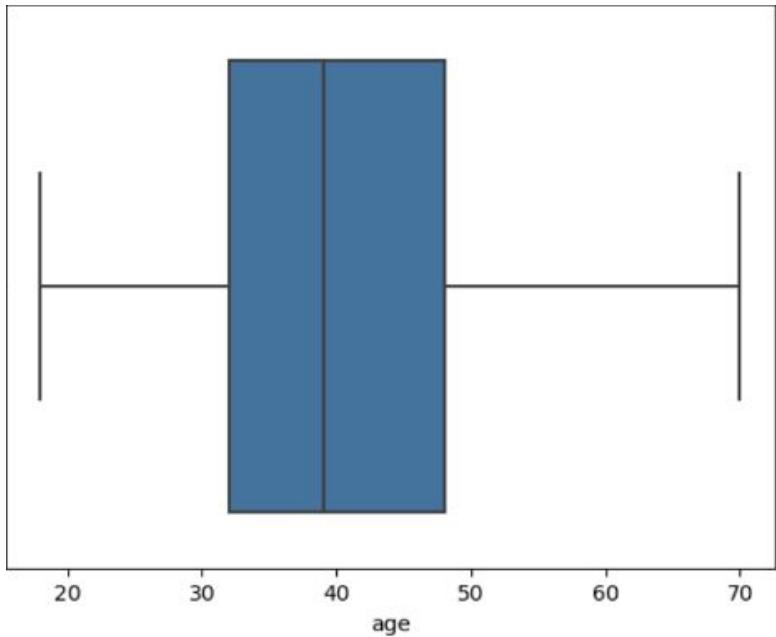
Removing outliers for “Age”



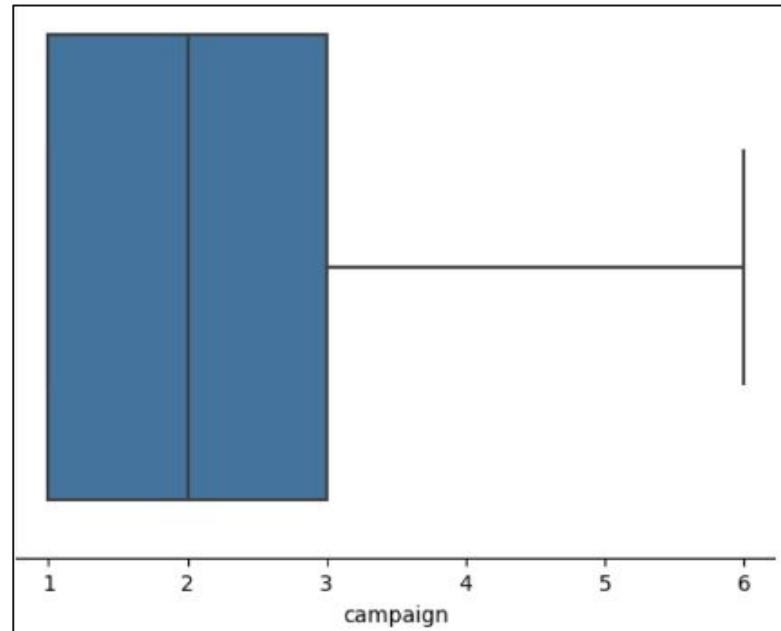
Removing outliers for “Campaign”

Data Preprocessing contd...

After Removing “Outliers”



Result of Removing outliers for “Age”



Result of Removing outliers for “Campaign”

Encoding & Standardization

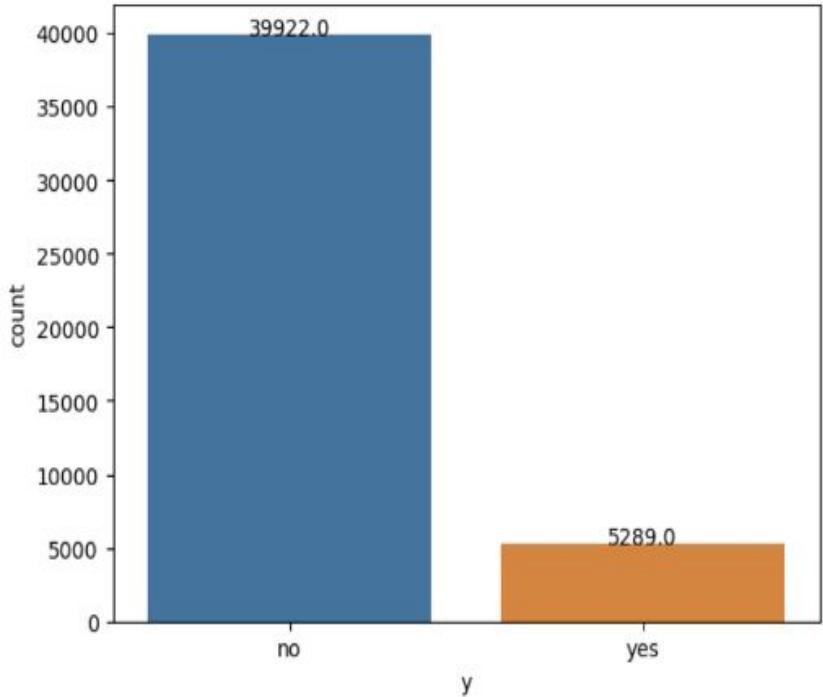
Data set on which we are encoding and standardizing

	age	job	marital	education	default	balance	housing	loan	duration	campaign	y
23787	47	blue-collar	married	primary	no	961	no	no	45	11	no
36427	37	technician	single	tertiary	no	555	yes	no	180	3	no
1547	53	blue-collar	married	secondary	no	359	no	yes	238	8	no
33652	29	blue-collar	married	secondary	no	457	yes	no	72	3	no
27706	30	blue-collar	single	tertiary	no	4127	yes	no	47	4	no

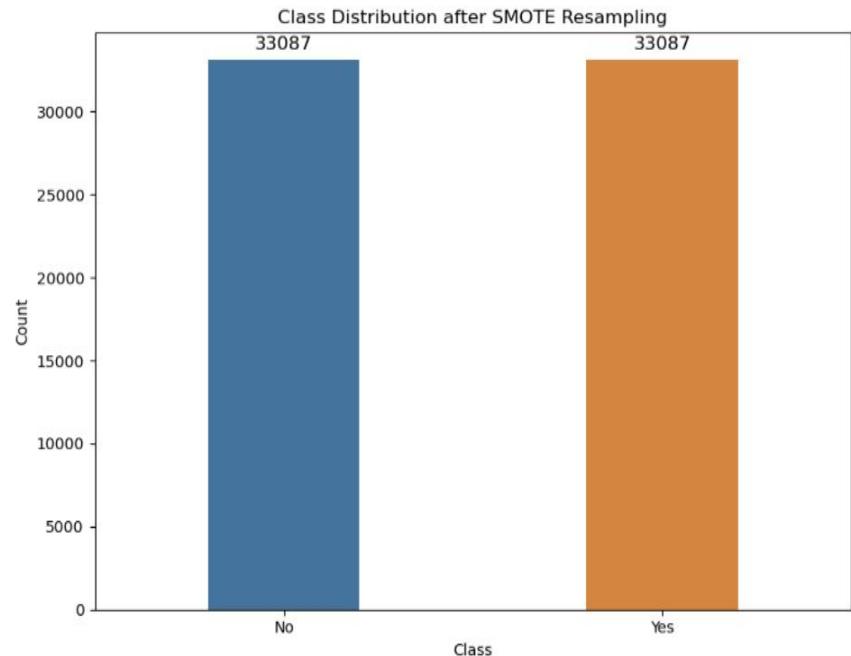
	age	job	marital	education	default	balance	housing	loan	duration	campaign	y
41748	0.269231	4	2	2	0	0.916055	1	0	0.069054	0.0	0
24558	0.653846	4	1	2	0	0.602861	1	0	0.029116	0.0	0
15520	0.269231	9	2	1	0	0.795219	1	0	0.062613	0.0	0
40919	0.538462	9	1	1	0	0.354602	0	0	0.027570	0.0	0
35534	0.365385	1	1	1	0	0.493883	1	0	0.111827	0.4	0

Output

SMOTE Sampling Technique



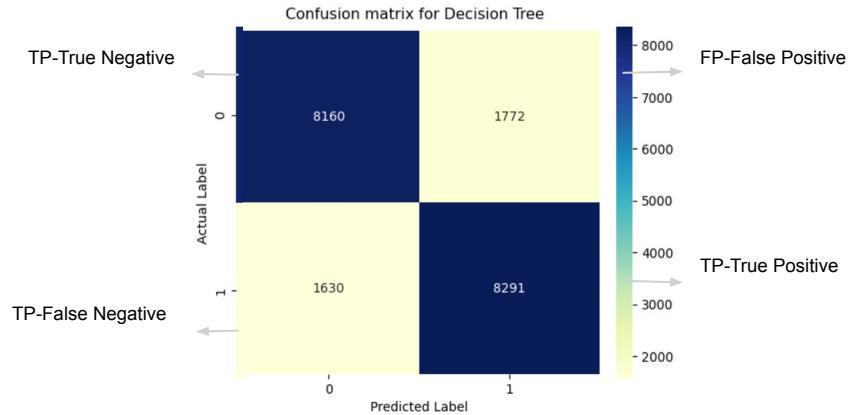
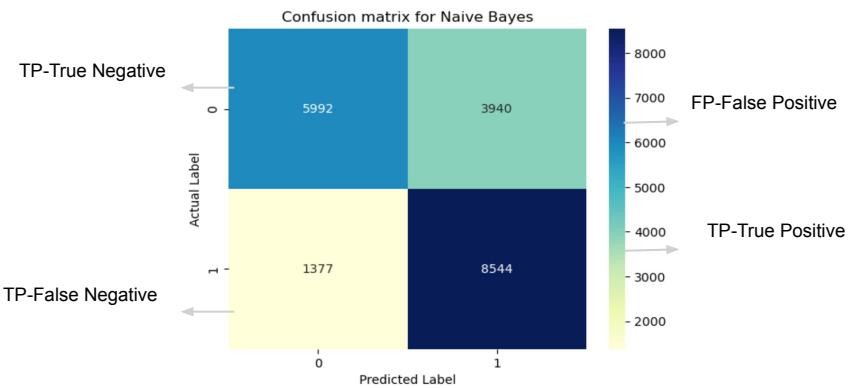
Pre-Sampling



Post-Sampling

Confusion Matrix(Accuracy, Precision, Recall, and F1 Score)+Optimal K-Value

Confusion Matrix



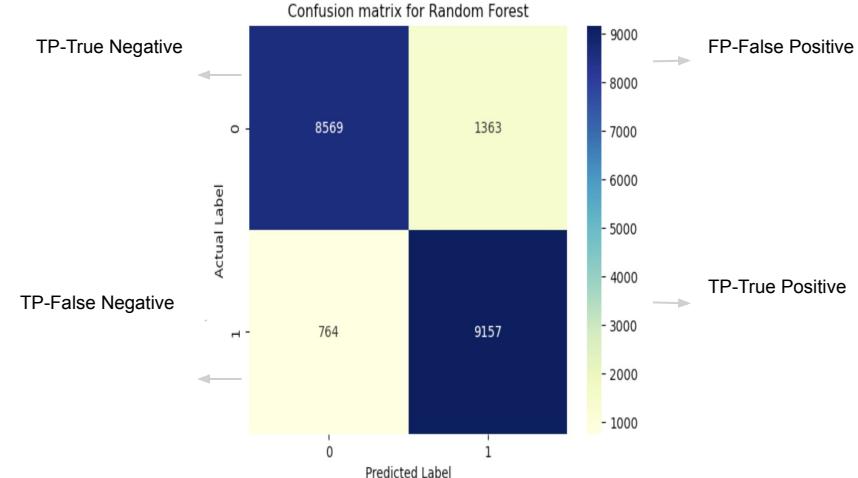
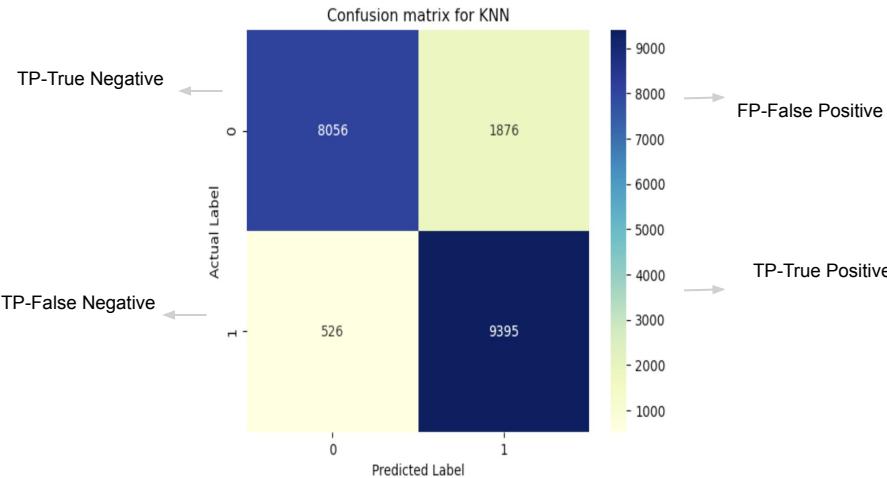
Model: Naive Bayes

Confusion Matrix: [[5992 3940], [1377 8544]]
Accuracy: 0.73
Precision: 0.68
PRecall: 0.86
F1 Score: 0.76

Model: Decision Tree

Confusion Matrix: [[8160 1772], [1630 8291]]
Accuracy: 0.83
Precision: 0.82
PRecall: 0.84
F1 Score: 0.83

Confusion Matrix



Model: KNN

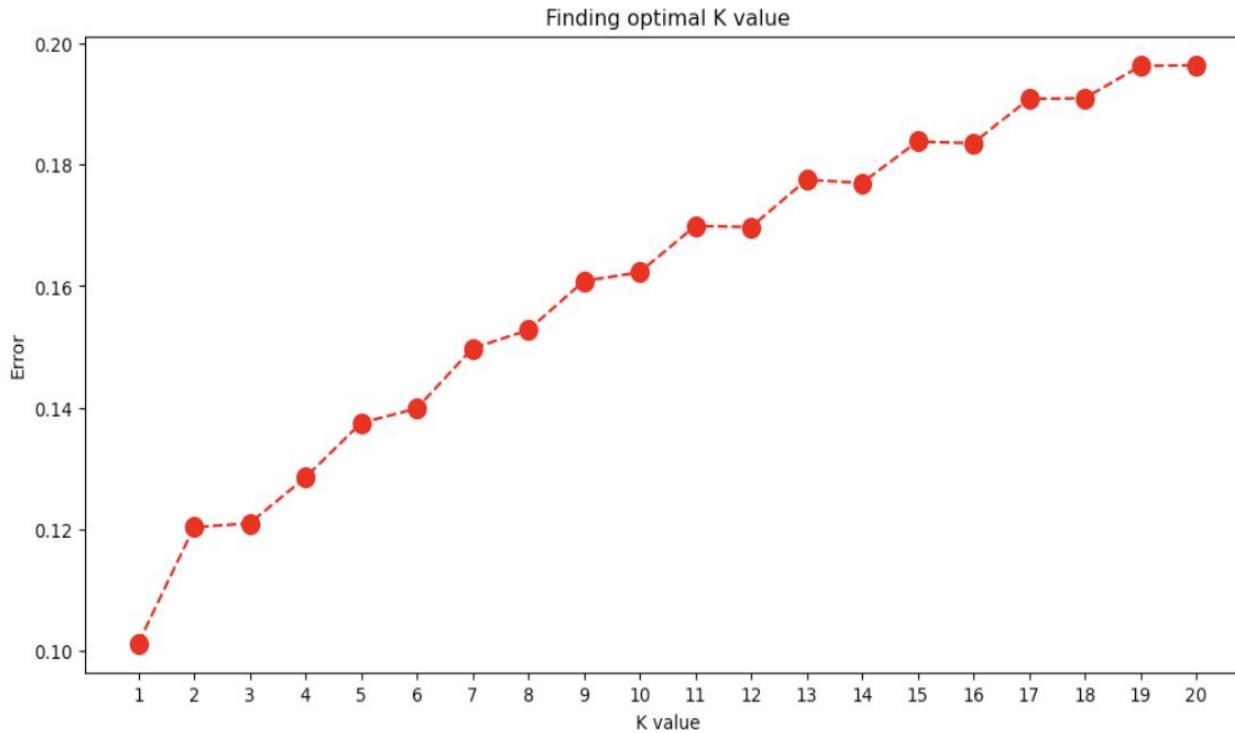
Confusion Matrix: [[8056 1876], [526 9395]]
Accuracy: 0.88
Precision: 0.83
PRecall: 0.95
F1 Score: 0.89



Model: Random Forest

Confusion Matrix: [[8569 1363], [764 9157]]
Accuracy: 0.89
Precision: 0.87
PRecall: 0.92
F1 Score: 0.90

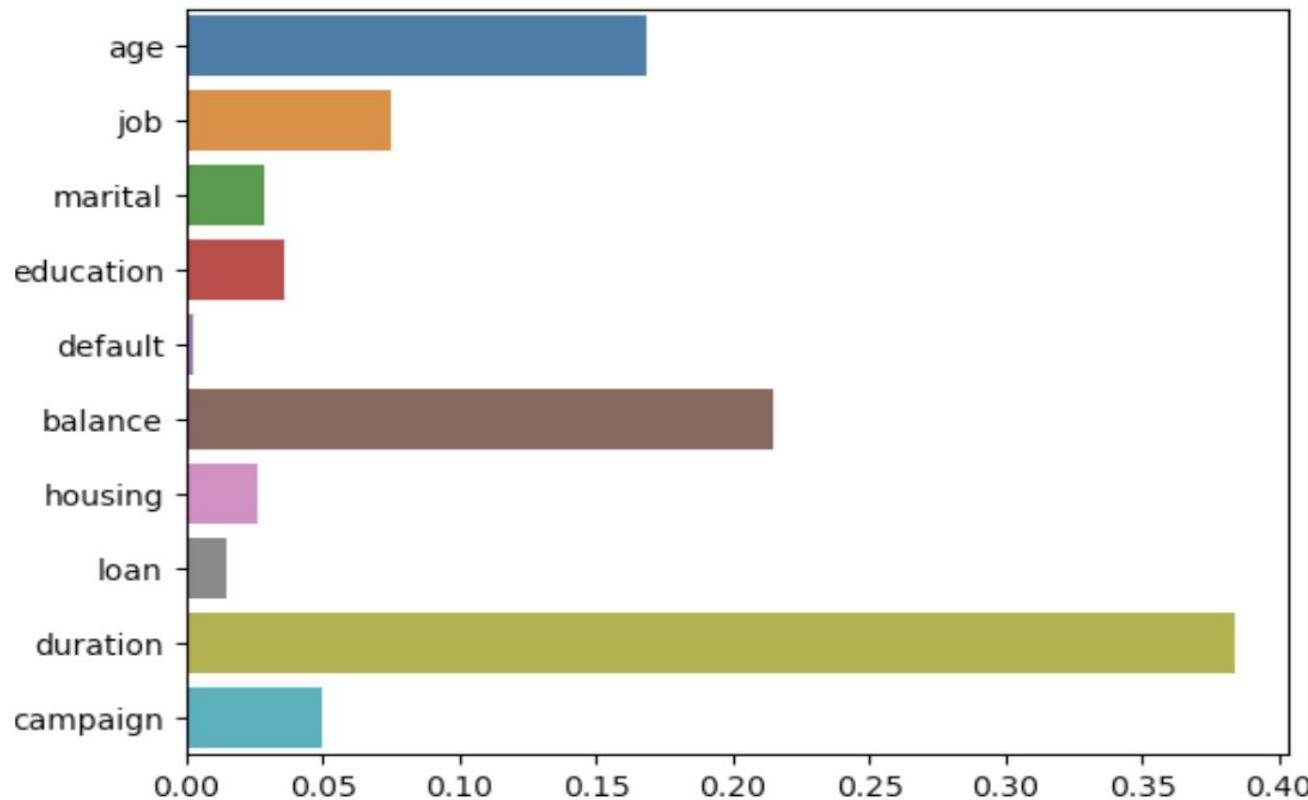
To find the Optimal K-Value



Optimal K-Value-3

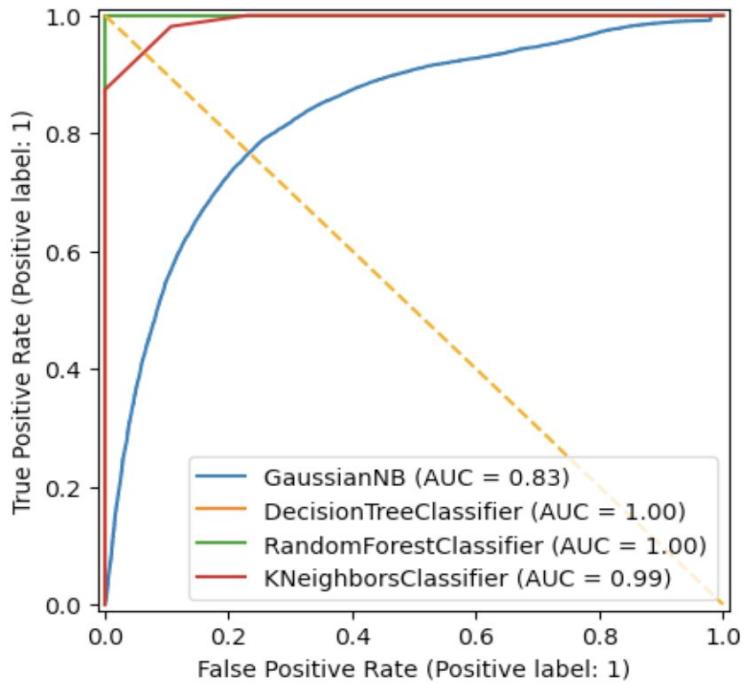
Feature Importance of Random Forest

Feature Importance for Random Forest

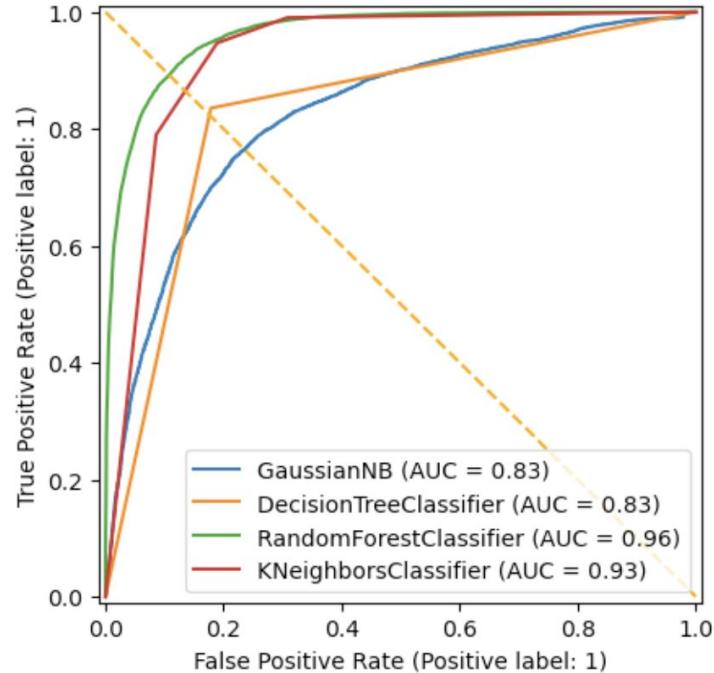


ROC

ROC



ROC-Training Dataset



ROC-Testing Dataset

Sample Data - Testing

Sample Data Testing

Sample Data Set

	age	job	marital	education	default	balance	housing	loan	duration	campaign
0	35	teacher	single	tertiary	no	1200	yes	no	320	3
1	28	artist	divorced	secondary	no	850	no	yes	145	2
2	52	entrepreneur	married	tertiary	yes	2100	yes	yes	180	1
3	43	engineer	married	secondary	no	3500	no	no	400	4

Testing Data Set - O/P

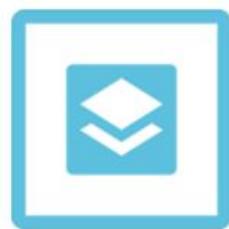
	age	job	marital	education	default	balance	housing	loan	duration	campaign	pred_NB	pred_DT	pred_KNN	pred_RF
0	35	teacher	single	tertiary	no	1200	yes	no	320	3	yes	no	no	yes
1	28	artist	divorced	secondary	no	850	no	yes	145	2	no	no	no	no
2	52	entrepreneur	married	tertiary	yes	2100	yes	yes	180	1	no	yes	no	yes
3	43	engineer	married	secondary	no	3500	no	no	400	4	yes	yes	no	yes

Regression

Dataset



Indian Data



House Price India



Dataset

- <https://data.world/dataindianset2000/house-price-india>

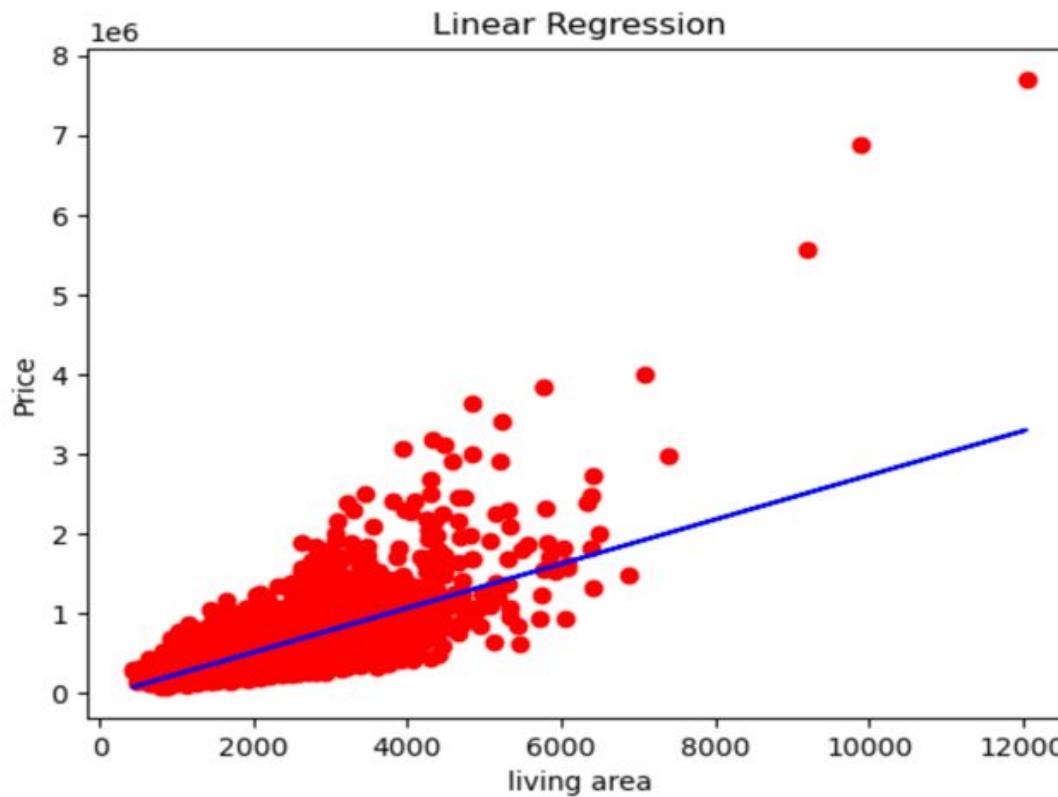
Dataset-Preprocessing

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   number of bedrooms      14620 non-null   int64  
 1   number of bathrooms     14620 non-null   float64 
 2   living area            14620 non-null   int64  
 3   lot area              14620 non-null   int64  
 4   number of floors       14620 non-null   float64 
 5   waterfront present    14620 non-null   int64  
 6   number of views        14620 non-null   int64  
 7   condition of the house 14620 non-null   int64  
 8   grade of the house     14620 non-null   int64  
 9   Area of the house(excluding basement) 14620 non-null   int64  
 10  Area of the basement   14620 non-null   int64  
 11  Built Year            14620 non-null   int64  
 12  living_area_renov     14620 non-null   int64  
 13  lot_area_renov        14620 non-null   int64  
 14  Number of schools nearby 14620 non-null   int64  
 15  Distance from the airport 14620 non-null   int64  
 16  Price                  14620 non-null   int64
```

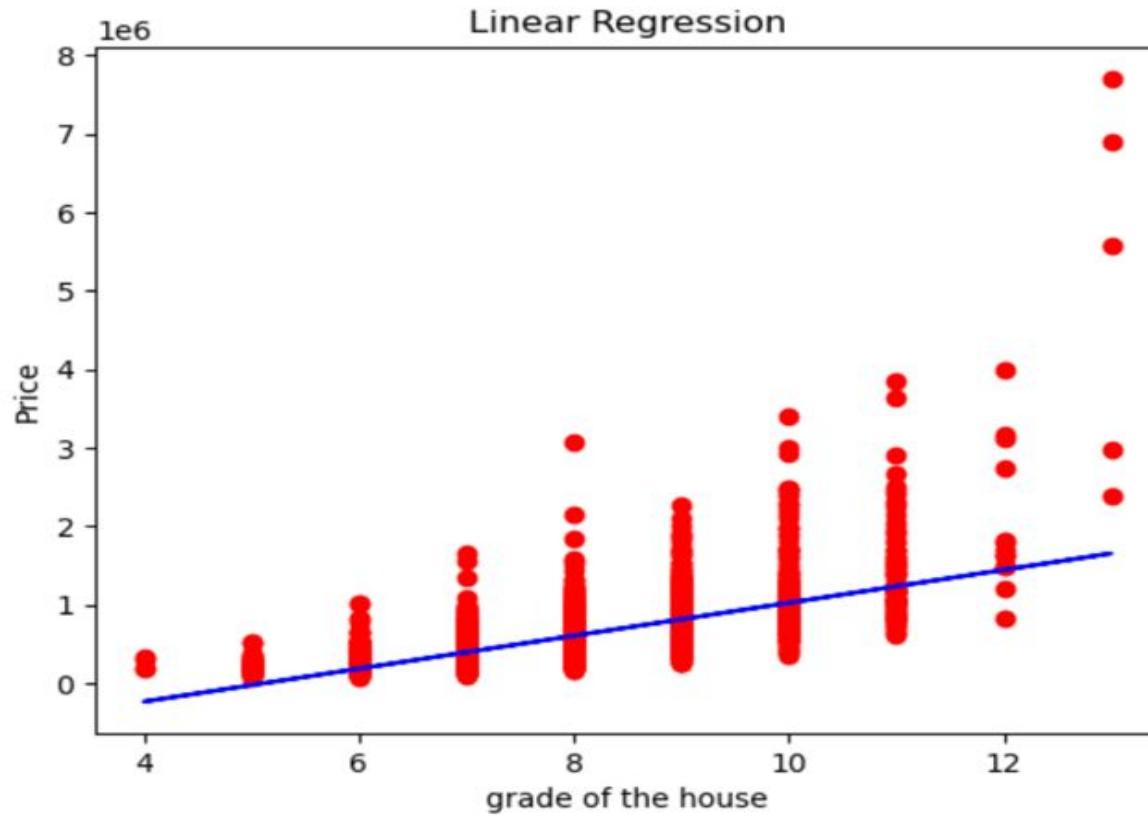
Correlation



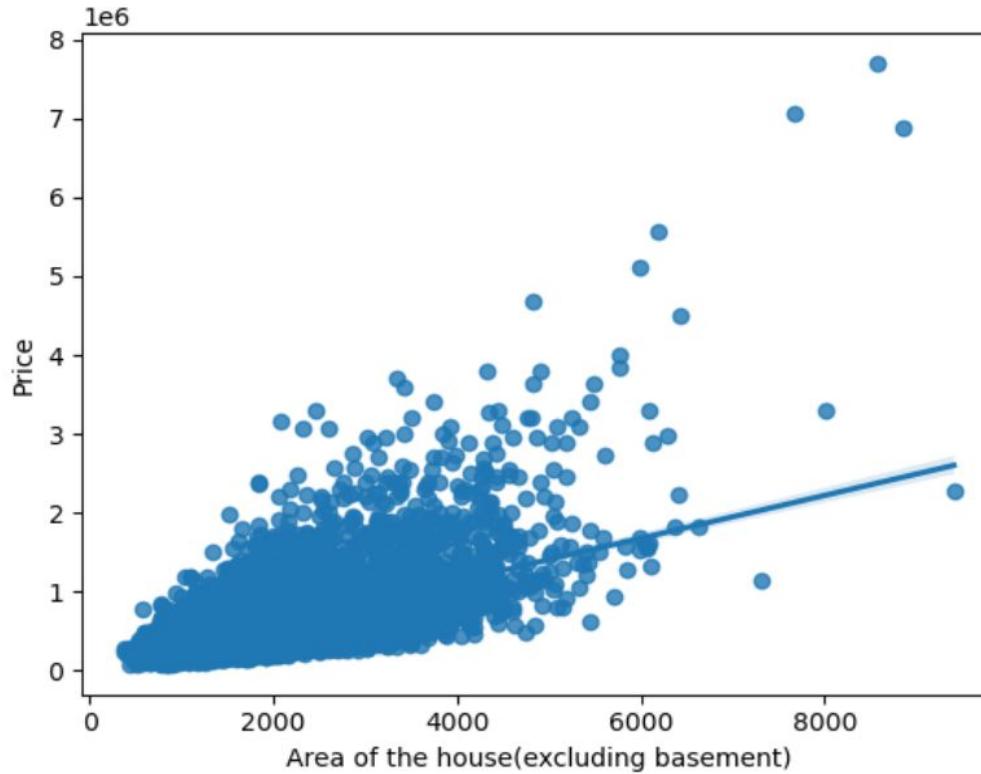
Regression plot(X = Living Area VS Y = Price)



Regression plot(X = Grade of House VS Y = Price)



Regression plot(X = Area of the House VS Y = Price)



Regression-Multiple Linear

Summary of Multiple Linear Regression

OLS Regression Results								
Dep. Variable:	Price	R-squared:	0.664					
Model:	OLS	Adj. R-squared:	0.664					
Method:	Least Squares	F-statistic:	1926.					
Date:	Wed, 02 Oct 2024	Prob (F-statistic):	0.00					
Time:	10:20:48	Log-Likelihood:	-2.0012e+05					
No. Observations:	14620	AIC:	4.003e+05					
Df Residuals:	14604	BIC:	4.004e+05					
Df Model:	15							
Covariance Type:	nonrobust							
		coef	std err	t	P> t	[0.025	0.975]	
const	✓	6.156e+06	1.57e+05	39.153	0.000	5.85e+06	6.46e+06	
number of bedrooms		-3.984e+04	2387.221	-16.688	0.000	-4.45e+04	-3.52e+04	
number of bathrooms		4.652e+04	4141.671	11.232	0.000	3.84e+04	5.46e+04	
living area		114.4892	2.876	39.814	0.000	108.853	120.126	
lot area		-0.1311	0.066	-1.985	0.047	-0.261	-0.002	
number of floors		2.474e+04	4536.891	5.454	0.000	1.59e+04	3.36e+04	
waterfront present		5.77e+05	2.21e+04	26.078	0.000	5.34e+05	6.2e+05	
number of views		4.217e+04	2736.267	15.410	0.000	3.68e+04	4.75e+04	
condition of the house		2.151e+04	2921.274	7.364	0.000	1.58e+04	2.72e+04	
grade of the house		1.189e+05	2712.828	43.834	0.000	1.14e+05	1.24e+05	
Area of the house(excluding basement)		58.5015	2.827	20.691	0.000	52.959	64.043	
Area of the basement		55.9878	3.313	16.901	0.000	49.495	62.481	
Built Year		-3550.9701	80.317	-44.212	0.000	-3708.402	-3393.538	
living_area_renov		13.0214	4.328	3.009	0.003	4.538	21.504	
lot_area_renov		-0.5166	0.097	-5.331	0.000	-0.707	-0.327	
Number of schools nearby		3040.0393	2157.430	1.409	0.159	-1188.796	7268.875	
Distance from the airport		-113.8137	197.351	-0.577	0.564	-500.647	273.019	
Omnibus:	11983.201	Durbin-Watson:	1.624					
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1201714.919					
Skew:	3.347	Prob(JB):	0.00					
Kurtosis:	46.908	Cond. No.	2.95e+17					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 3.76e-22. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

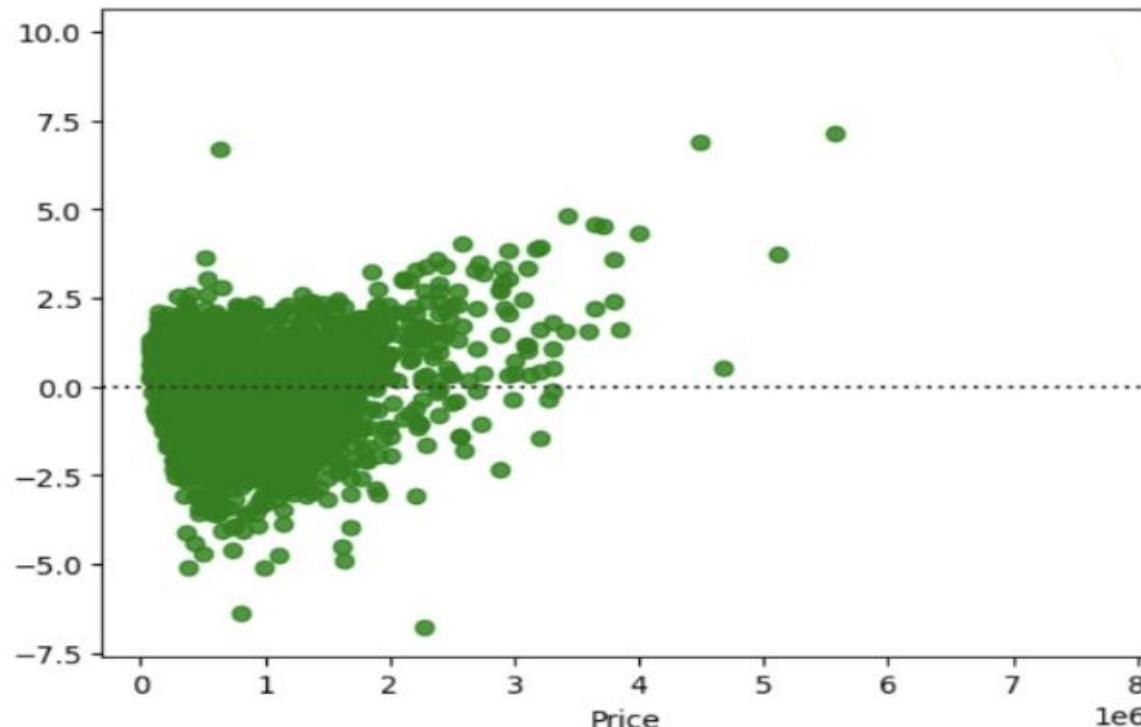
Error Matrix for Linear Regression

```
res= stat()  
res.reg_metric(y=np.array(y), yhat=np.array(reg.predict(X)), resid=np.array(reg.resid))  
res.reg_metric_df
```

	Metrics	Value
0	Root Mean Square Error (RMSE)	2.129635e+05
1	Mean Squared Error (MSE)	4.535346e+10
2	Mean Absolute Error (MAE)	1.366337e+05
3	Mean Absolute Percentage Error (MAPE)	2.851000e-01

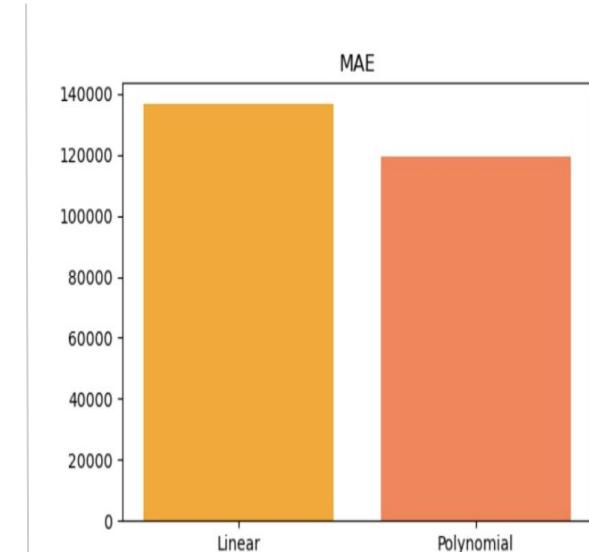
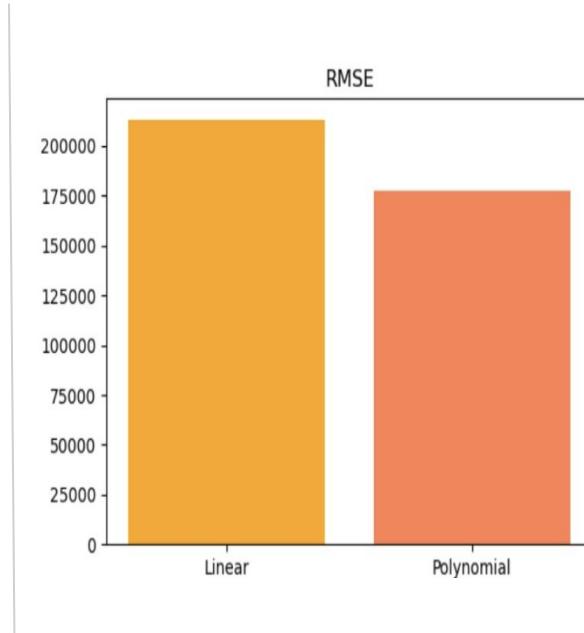
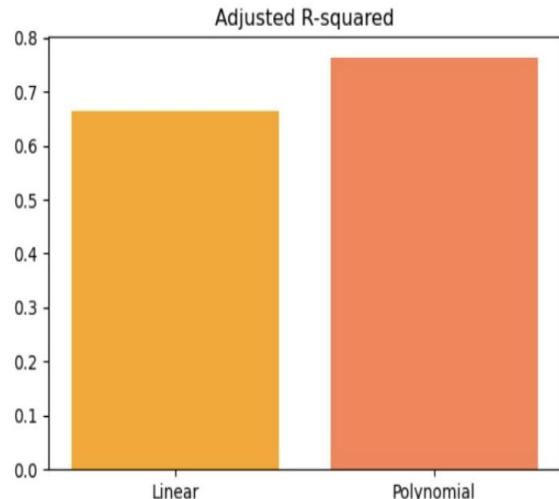
Residual plot

```
influence=reg.get_influence()  
residual_price=influence.resid_studentized_external  
sns.residplot(x = 'Price',y = residual_price,data = df,color='g');
```



Regression-Polynomial

Comparing Multiple Linear and Polynomial Regression



Polynomial regression is the best regression model for the chosen dataset because the **AR-squared value is higher** for Poly-regression VS Linear and **RMSE and MAE are lower**.

Thank You!

scm516_group_project_final

October 5, 2024

[3]: # Necessary imports for preprocessing and machine learning

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer    # interpolation for missing values
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import RocCurveDisplay, confusion_matrix, accuracy_score,precision_score, recall_score, f1_score, classification_report
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
from bioinfokit.analys import stat
from bioinfokit import visuz
from sklearn.preprocessing import PolynomialFeatures
```

[28]: # load the dataset

```
df_bank=pd.read_csv("bank-full.csv", delimiter=';')
```

[29]: df_bank

```
[29]:      age         job marital education default  balance housing loan \
0       58   management  married   tertiary     no     2143    yes    no
1       44  technician   single  secondary     no      29    yes    no
2       33  entrepreneur  married  secondary     no       2    yes   yes
3       47  blue-collar  married   unknown     no    1506    yes    no
4       33        unknown   single   unknown     no       1    no    no
...     ...        ...     ...     ...     ...     ...     ...     ...     ...
```

```

45206 51 technician married tertiary no 825 no no
45207 71 retired divorced primary no 1729 no no
45208 72 retired married secondary no 5715 no no
45209 57 blue-collar married secondary no 668 no no
45210 37 entrepreneur married secondary no 2971 no no

      contact day month duration campaign pdays previous poutcome y
0 unknown 5 may 261 1 -1 0 unknown no
1 unknown 5 may 151 1 -1 0 unknown no
2 unknown 5 may 76 1 -1 0 unknown no
3 unknown 5 may 92 1 -1 0 unknown no
4 unknown 5 may 198 1 -1 0 unknown no
...
45206 cellular 17 nov 977 3 -1 0 unknown yes
45207 cellular 17 nov 456 2 -1 0 unknown yes
45208 cellular 17 nov 1127 5 184 3 success yes
45209 telephone 17 nov 508 4 -1 0 unknown no
45210 cellular 17 nov 361 2 188 11 other no

```

[45211 rows x 17 columns]

[30]: # Descriptive statistics of data

```
df_bank.describe()
```

	age	balance	day	duration	campaign
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.936210	1362.272058	15.806419	258.163080	2.763841
std	10.618762	3044.765829	8.322476	257.527812	3.098021
min	18.000000	-8019.000000	1.000000	0.000000	1.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000
	pdays	previous			
count	45211.000000	45211.000000			
mean	40.197828	0.580323			
std	100.128746	2.303441			
min	-1.000000	0.000000			
25%	-1.000000	0.000000			
50%	-1.000000	0.000000			
75%	-1.000000	0.000000			
max	871.000000	275.000000			

[31]: df_bank.drop(columns='contact', inplace=True)
df_bank.drop(columns='previous', inplace=True)

```
df_bank.drop(columns='poutcome', inplace=True)
df_bank.drop(columns='pdays', inplace=True)
df_bank.drop(columns='month', inplace=True)
df_bank.drop(columns='day', inplace=True)
```

[32]: df_bank

```
[32]:      age         job marital education default balance housing loan \
0       58 management married   tertiary    no     2143    yes   no
1       44 technician single secondary   no      29    yes   no
2       33 entrepreneur married secondary   no       2    yes   yes
3       47 blue-collar married unknown    no    1506    yes   no
4       33          unknown single unknown    no       1    no   no
...
45206    51 technician married tertiary    no     825    no   no
45207    71      retired divorced primary   no    1729    no   no
45208    72      retired married secondary   no    5715    no   no
45209    57 blue-collar married secondary   no     668    no   no
45210    37 entrepreneur married secondary   no    2971    no   no

      duration campaign    y
0        261         1  no
1        151         1  no
2         76         1  no
3         92         1  no
4        198         1  no
...
45206     977         3 yes
45207     456         2 yes
45208    1127         5 yes
45209     508         4 no
45210     361         2 no

[45211 rows x 11 columns]
```

[33]: # check the type of data and see if there is any missing value

```
df_bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 11 columns):
 #   Column      Non-Null Count Dtype
 ---  --          --          --
 0   age          45211 non-null int64
 1   job           45211 non-null object
 2   marital       45211 non-null object
 3   education     45211 non-null object
```

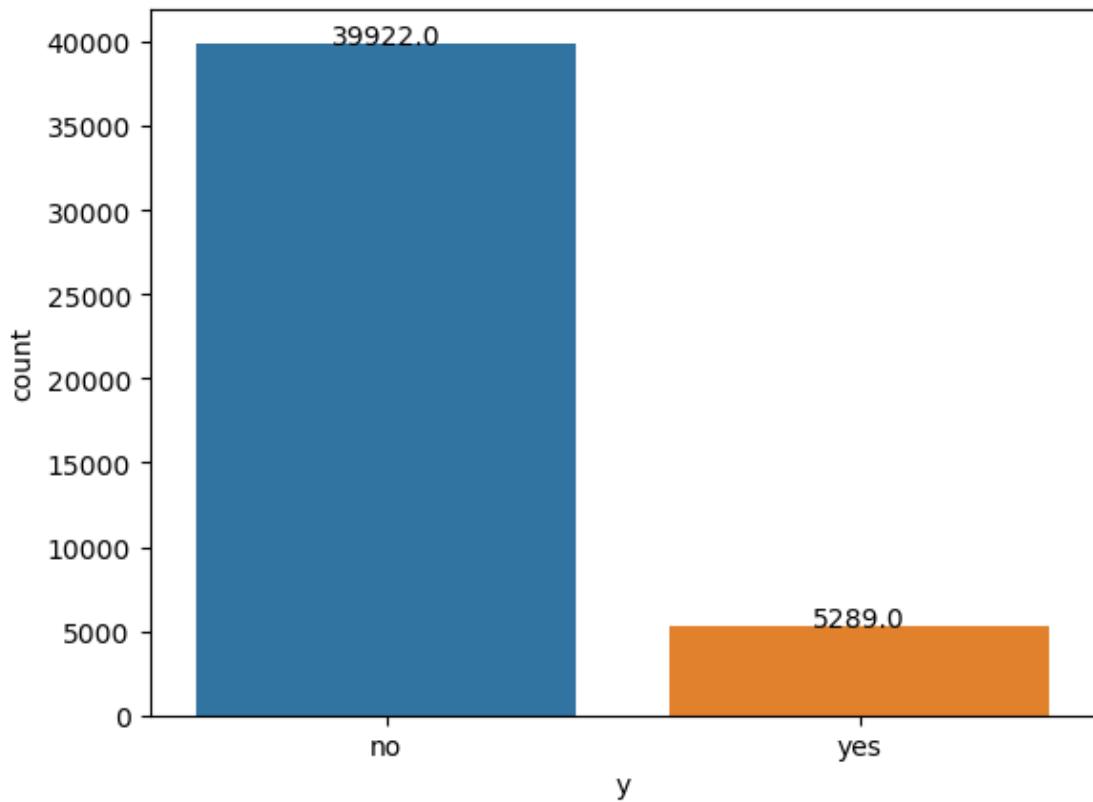
```
4    default    45211 non-null  object
5    balance    45211 non-null  int64
6    housing    45211 non-null  object
7    loan        45211 non-null  object
8    duration   45211 non-null  int64
9    campaign   45211 non-null  int64
10   y          45211 non-null  object
dtypes: int64(4), object(7)
memory usage: 3.8+ MB
```

```
[34]: # Is the dataset balanced?
```

```
# Create the countplot
ax = sns.countplot(data=df_bank, x='y')

# Add count labels on top of the bars
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
                                         get_height(),
                                         ha='center', va='baseline')

# Display the plot
plt.show()
```



```
[35]: # Number of people with term_deposit  
(df_bank[['y']]=="yes").sum()
```

```
[35]: y      5289  
dtype: int64
```

```
[36]: # Number of people without term_deposit  
(df_bank[['y']]=="no").sum()
```

```
[36]: y      39922  
dtype: int64
```

```
[37]: # calculate probability of taking term_deposit  
(df_bank[['y']]=="yes").sum() / len(df_bank['y'])
```

```
[37]: y      0.116985  
dtype: float64
```

```
[38]: # How to get the list of all columns in my dataset  
  
columns_list=df_bank.columns.tolist()
```

```
[39]: columns_list
```

```
[39]: ['age',  
       'job',  
       'marital',  
       'education',  
       'default',  
       'balance',  
       'housing',  
       'loan',  
       'duration',  
       'campaign',  
       'y']
```

```
[40]: # We need to create a variable that stores the name of categorical columns  
# as well a varibale that stores the name of numerical variable.  
# because we need to handle missing values for categorical using mode technique  
# we also need to replace the missing values of numerical variable using mean.  
# You can also use linear interpolation  
  
# List of categorical variables (dtype == object or category)
```

```

categorical_cols = df_bank.select_dtypes(include=['object', 'category']).columns.tolist()

# List of numerical variables (dtype == int or float)
numerical_cols = df_bank.select_dtypes(include=['number']).columns.tolist()

# Display the lists
print("Categorical columns:", categorical_cols)
print("Numerical columns:", numerical_cols)

```

Categorical columns: ['job', 'marital', 'education', 'default', 'housing',
'loan', 'y']
Numerical columns: ['age', 'balance', 'duration', 'campaign']

[41]: df_bank.isnull().sum()

```

[41]: age      0
       job     0
       marital 0
       education 0
       default 0
       balance 0
       housing 0
       loan    0
       duration 0
       campaign 0
       y        0
       dtype: int64

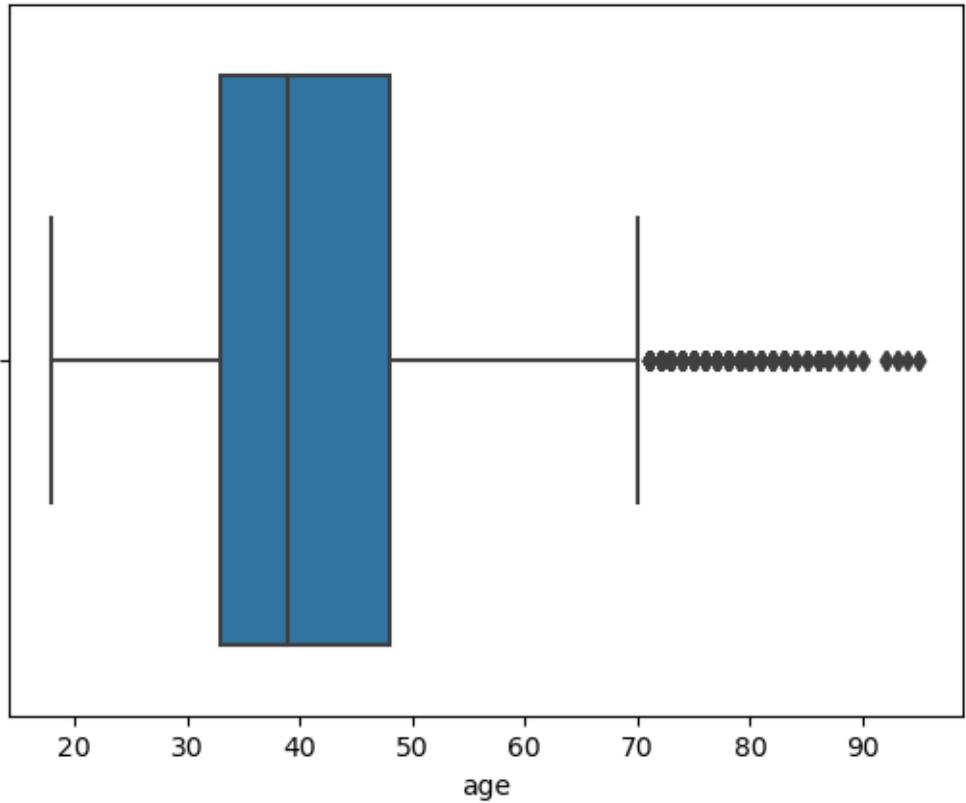
```

[42]: # Find and remove outlier from numerical variable

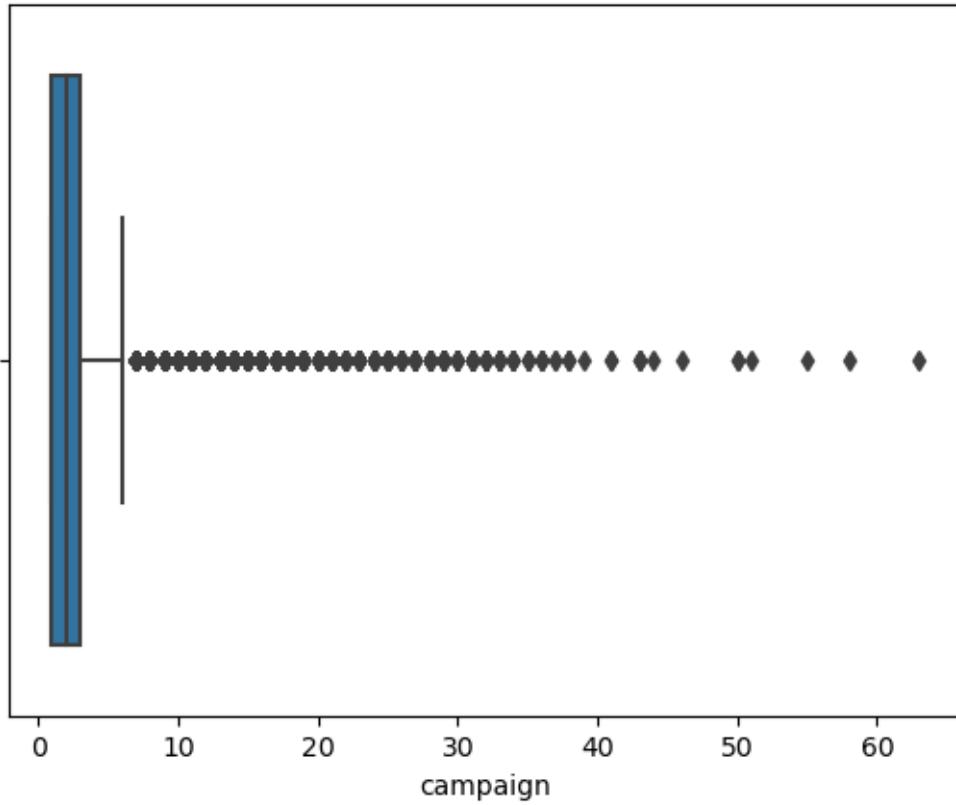
```

# Create a boxplot for the 'age' column to detect outliers
sns.boxplot(data=df_bank, x='age');      # there are outliers (black dots). Those
                                         # outliers should be removed using IQR method

```



```
[43]: sns.boxplot(data=df_bank, x='campaign');
```



```
[44]: # Calculate the IQR for the 'campaign' column
Q1 = df_bank['campaign'].quantile(0.25) # 25th percentile
Q3 = df_bank['campaign'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1 # Interquartile range
# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
# Filter out the outliers
df_bank = df_bank[(df_bank['campaign'] >= lower_bound) & (df_bank['campaign'] <= upper_bound)]
# df1 is a new datafram after removing outliers
# Display the results (rows without outliers)
```

```
[45]: df_bank
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	

```

...
45206 51 technician married tertiary no 825 no no
45207 71 retired divorced primary no 1729 no no
45208 72 retired married secondary no 5715 no no
45209 57 blue-collar married secondary no 668 no no
45210 37 entrepreneur married secondary no 2971 no no

duration campaign y
0 261 1 no
1 151 1 no
2 76 1 no
3 92 1 no
4 198 1 no
...
45206 977 3 yes
45207 456 2 yes
45208 1127 5 yes
45209 508 4 no
45210 361 2 no

```

[42147 rows x 11 columns]

```
[46]: # Calculate the IQR for the 'age' column
Q1 = df_bank['age'].quantile(0.25) # 25th percentile
Q3 = df_bank['age'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1 # Interquartile range
# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
# Filter out the outliers
df_bank = df_bank[(df_bank['age'] >= lower_bound) & (df_bank['age'] <=
    ↪upper_bound)]
```

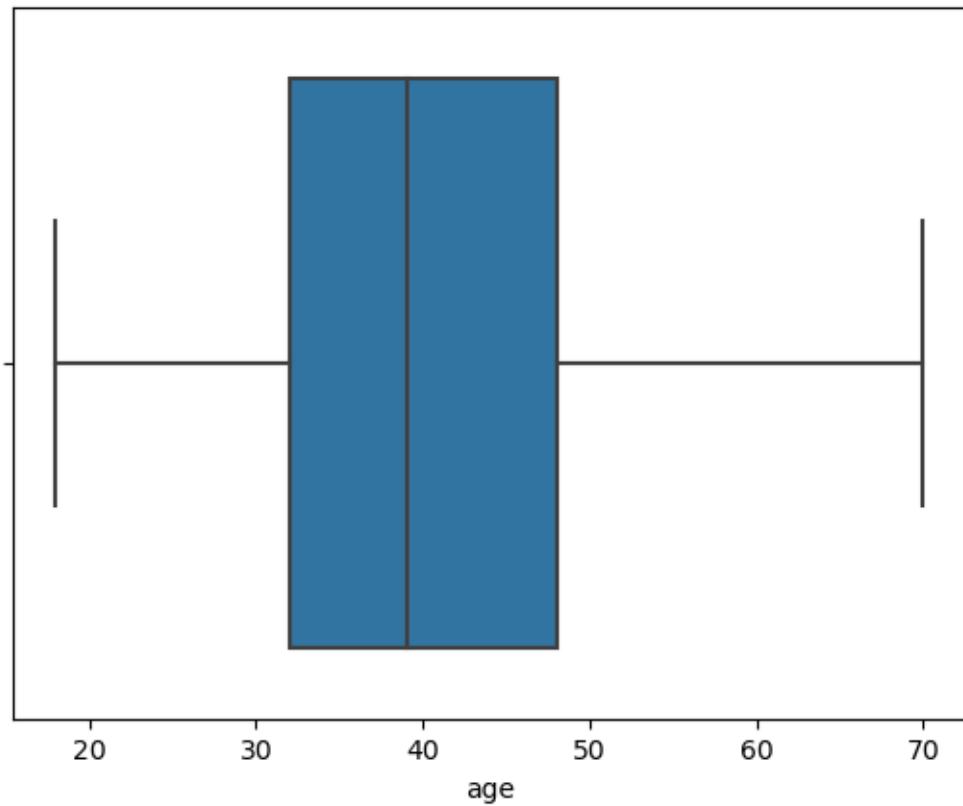
[47]: df_bank

```
[47]: age job marital education default balance housing loan \
0 58 management married tertiary no 2143 yes no
1 44 technician single secondary no 29 yes no
2 33 entrepreneur married secondary no 2 yes yes
3 47 blue-collar married unknown no 1506 yes no
4 33 unknown single unknown no 1 no no
...
45203 23 student single tertiary no 113 no no
45205 25 technician single secondary no 505 no yes
45206 51 technician married tertiary no 825 no no
45209 57 blue-collar married secondary no 668 no no
45210 37 entrepreneur married secondary no 2971 no no
```

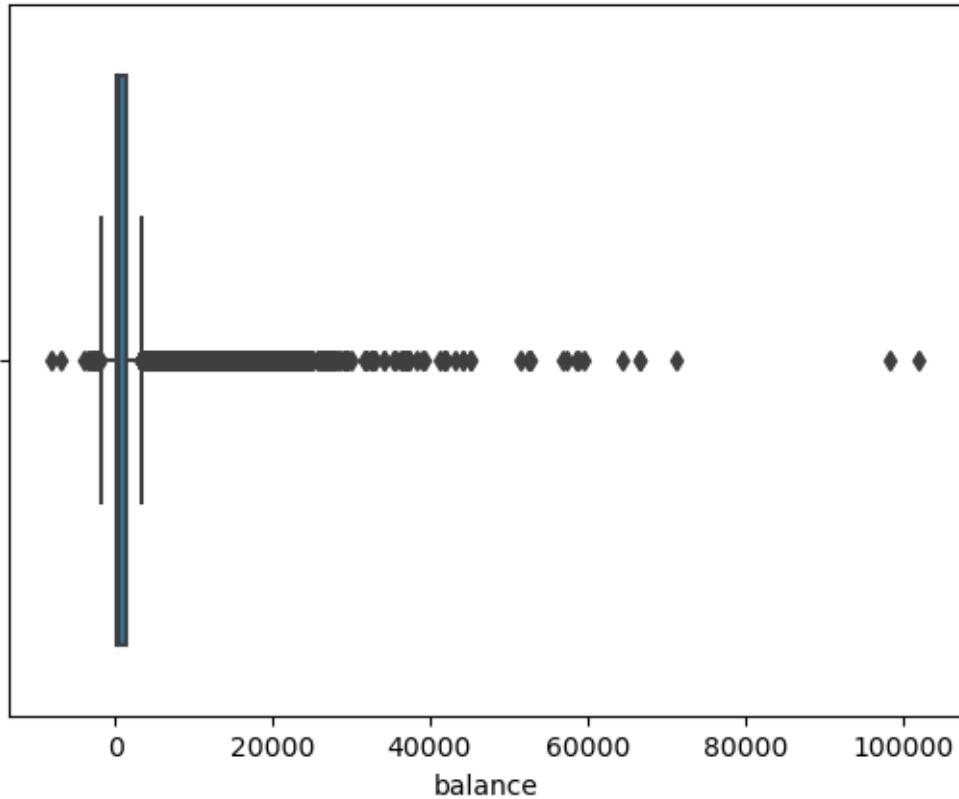
```
duration campaign y
0          261      1 no
1          151      1 no
2           76      1 no
3           92      1 no
4          198      1 no
...
45203       266      1 yes
45205       386      2 yes
45206       977      3 yes
45209       508      4 no
45210       361      2 no
```

[41674 rows x 11 columns]

```
[48]: sns.boxplot(data=df_bank, x='age');
```



```
[49]: sns.boxplot(data=df_bank, x='balance');
```



```
[50]: # Calculate the IQR for the 'balance' column
Q1 = df_bank['balance'].quantile(0.25) # 25th percentile
Q3 = df_bank['balance'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1 # Interquartile range
# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
# Filter out the outliers
df_bank = df_bank[(df_bank['balance'] >= lower_bound) & (df_bank['balance'] <= upper_bound)]
df_bank
```

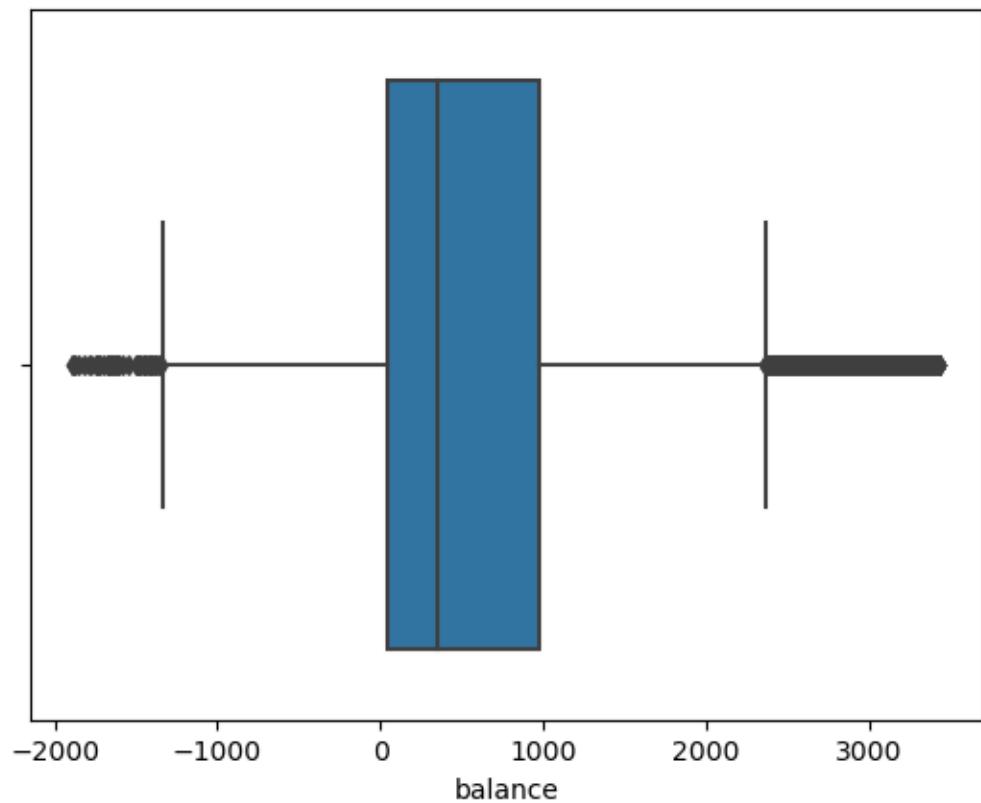
	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...	
45203	23	student	single	tertiary	no	113	no	no	
45205	25	technician	single	secondary	no	505	no	yes	

```
45206    51    technician married tertiary      no     825      no  no
45209    57    blue-collar married secondary    no     668      no  no
45210    37 entrepreneur married secondary    no    2971      no  no
```

```
duration campaign    y
0           261       1  no
1           151       1  no
2            76       1  no
3            92       1  no
4           198       1  no
...
45203      266       1 yes
45205      386       2 yes
45206      977       3 yes
45209      508       4  no
45210      361       2  no
```

[37303 rows x 11 columns]

```
[51]: sns.boxplot(data=df_bank, x='balance');
```



```
[52]: df_bank
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...
45203	23	student	single	tertiary	no	113	no	no	
45205	25	technician	single	secondary	no	505	no	yes	
45206	51	technician	married	tertiary	no	825	no	no	
45209	57	blue-collar	married	secondary	no	668	no	no	
45210	37	entrepreneur	married	secondary	no	2971	no	no	
	duration	campaign	y						
0	261	1	no						
1	151	1	no						
2	76	1	no						
3	92	1	no						
4	198	1	no						
...						
45203	266	1	yes						
45205	386	2	yes						
45206	977	3	yes						
45209	508	4	no						
45210	361	2	no						

[37303 rows x 11 columns]

```
[53]: # label Encoding
# we use the following function to encode categorical variable

from sklearn.preprocessing import LabelEncoder

lb=LabelEncoder()

for col in categorical_cols:    # we define categorical_cols in line [73]
    df_bank.loc[:,col]=lb.fit_transform(df_bank[col])  # note that df1 is new
    ↵dataframe without outliers
```

```
[54]: df_bank.head(5)
```

	age	job	marital	education	default	balance	housing	loan	duration	\
0	58	4	1	2	0	2143	1	0	261	
1	44	9	2	1	0	29	1	0	151	
2	33	2	1	1	0	2	1	1	76	

3	47	1	1	3	0	1506	1	0	92
4	33	11	2	3	0	1	0	0	198

	campaign	y
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

[55]: # Standardize the numerical variable (Age variable)

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

df_bank.loc[:,numerical_cols] = scaler.fit_transform(df_bank[numerical_cols]) ↴
    ↪# we define list of numerical variable in line [73]
```

[56]: df_bank

	age	job	marital	education	default	balance	housing	loan	\
0	0.769231	4	1	2	0	0.757952	1	0	
1	0.500000	9	2	1	0	0.360060	1	0	
2	0.288462	2	1	1	0	0.354978	1	1	
3	0.557692	1	1	3	0	0.638058	1	0	
4	0.288462	11	2	3	0	0.354790	0	0	
...	
45203	0.096154	8	2	2	0	0.375871	0	0	
45205	0.134615	9	2	1	0	0.449652	0	1	
45206	0.634615	9	1	2	0	0.509881	0	0	
45209	0.750000	1	1	1	0	0.480331	0	0	
45210	0.365385	2	1	1	0	0.913796	0	0	
	duration	campaign	y						
0	0.067251	0.0	0						
1	0.038907	0.0	0						
2	0.019583	0.0	0						
3	0.023705	0.0	0						
4	0.051018	0.0	0						
...						
45203	0.068539	0.0	1						
45205	0.099459	0.2	1						
45206	0.251739	0.4	1						
45209	0.130894	0.6	0						
45210	0.093017	0.2	0						

```
[37303 rows x 11 columns]
```

```
[57]: # define features (X) and response/label (y)

X= df_bank.drop(columns=['y'])
y=df_bank['y']
```

```
[39]: # We saw earlier that the data is not balanced. We apply SMOTE sampling
      ↵technique to our imbalanced data
smote = SMOTE()
X_res, y_res = smote.fit_resample(X,y.astype(int))
```

```
[40]: y_res
```

```
[40]: 0      0
1      0
2      0
3      0
4      0
..
66169    1
66170    1
66171    1
66172    1
66173    1
Name: y, Length: 66174, dtype: int32
```

```
[41]: y_res.value_counts()  # the response/ label has been balanced
```

```
[41]: y
0      33087
1      33087
Name: count, dtype: int64
```

```
[42]: # split into training and testing test
```

```
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.
      ↵3, random_state=8)
```

```
[43]: # Naive Bayes Model
```

```
nb_model = GaussianNB()
nb_model.fit(X_train, y_train)
y_pred_nb = nb_model.predict(X_test)
```

```
[44]: # Decision Tree Model
```

```
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
```

[45]: # Random Forest Model

```
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
y_pred_rf=rf_model.predict(X_test)
```

[46]: # KNN model

```
knn_model =KNeighborsClassifier(n_neighbors=3)
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)
```

[47]: # Evaluation: Confusion Matrix and Metrics

```
def evaluate_model(y_test, y_pred, model_name):
    print(f"Model: {model_name}")
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
    print(f"Precision: {precision_score(y_test, y_pred):.2f}")
    print(f"Recall: {recall_score(y_test, y_pred):.2f}")
    print(f"F1 Score: {f1_score(y_test, y_pred):.2f}\n")
```

[48]: # Evaluate all models

```
evaluate_model(y_test, y_pred_nb, "Naive Bayes")
evaluate_model(y_test, y_pred_dt, "Decision Tree")
evaluate_model(y_test, y_pred_rf, "Random Forest")
evaluate_model(y_test, y_pred_knn, "KNN")
```

Model: Naive Bayes
Confusion Matrix:
[[5992 3940]
 [1377 8544]]
Accuracy: 0.73
Precision: 0.68
Recall: 0.86
F1 Score: 0.76

Model: Decision Tree
Confusion Matrix:
[[8160 1772]
 [1630 8291]]
Accuracy: 0.83

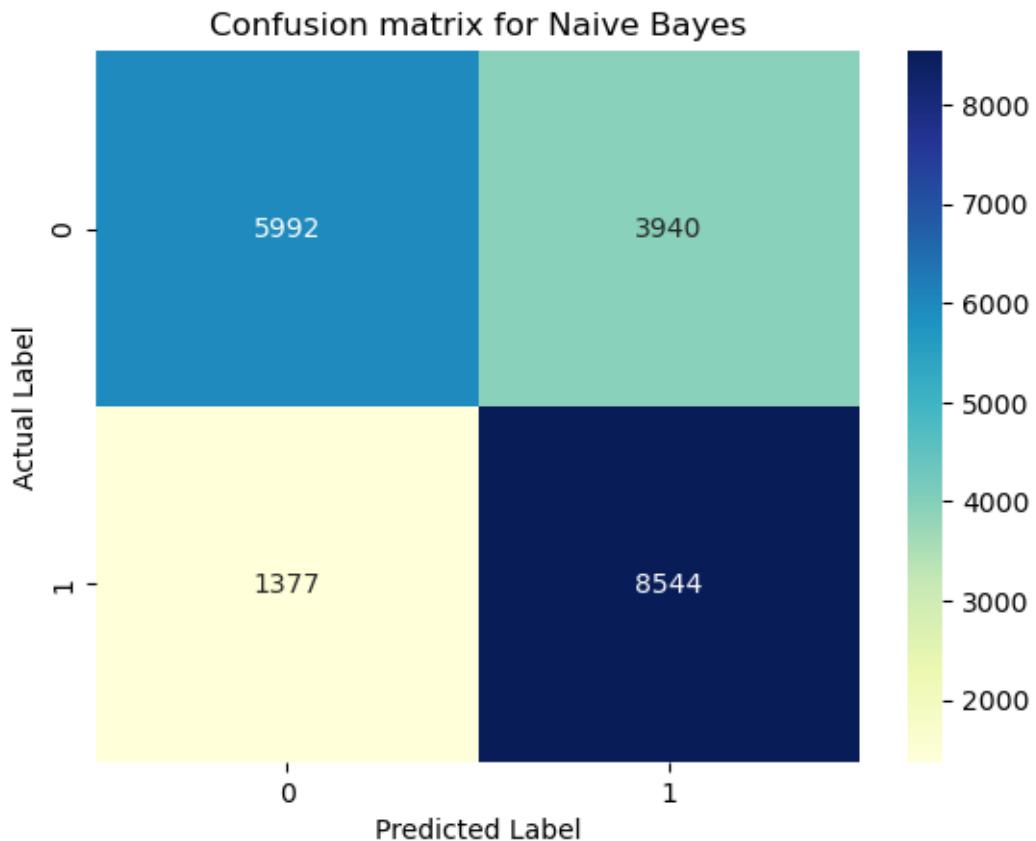
```
Precision: 0.82  
PRecall: 0.84  
F1 Score: 0.83
```

```
Model: Random Forest  
Confusion Matrix:  
[[8569 1363]  
 [ 764 9157]]  
Accuracy: 0.89  
Precision: 0.87  
PRecall: 0.92  
F1 Score: 0.90
```

```
Model: KNN  
Confusion Matrix:  
[[8056 1876]  
 [ 526 9395]]  
Accuracy: 0.88  
Precision: 0.83  
PRecall: 0.95  
F1 Score: 0.89
```

```
[49]: cm1=confusion_matrix(y_test, y_pred_nb)  
cm2=confusion_matrix(y_test, y_pred_dt)  
cm3=confusion_matrix(y_test, y_pred_rf)  
cm4=confusion_matrix(y_test, y_pred_knn)
```

```
[50]: # confusion matrix for Naive Bayes  
  
sns.heatmap(cm1, annot=True, cmap='YlGnBu', fmt='g')  
plt.ylabel('Actual Label')  
plt.xlabel('Predicted Label')  
plt.title('Confusion matrix for Naive Bayes');
```

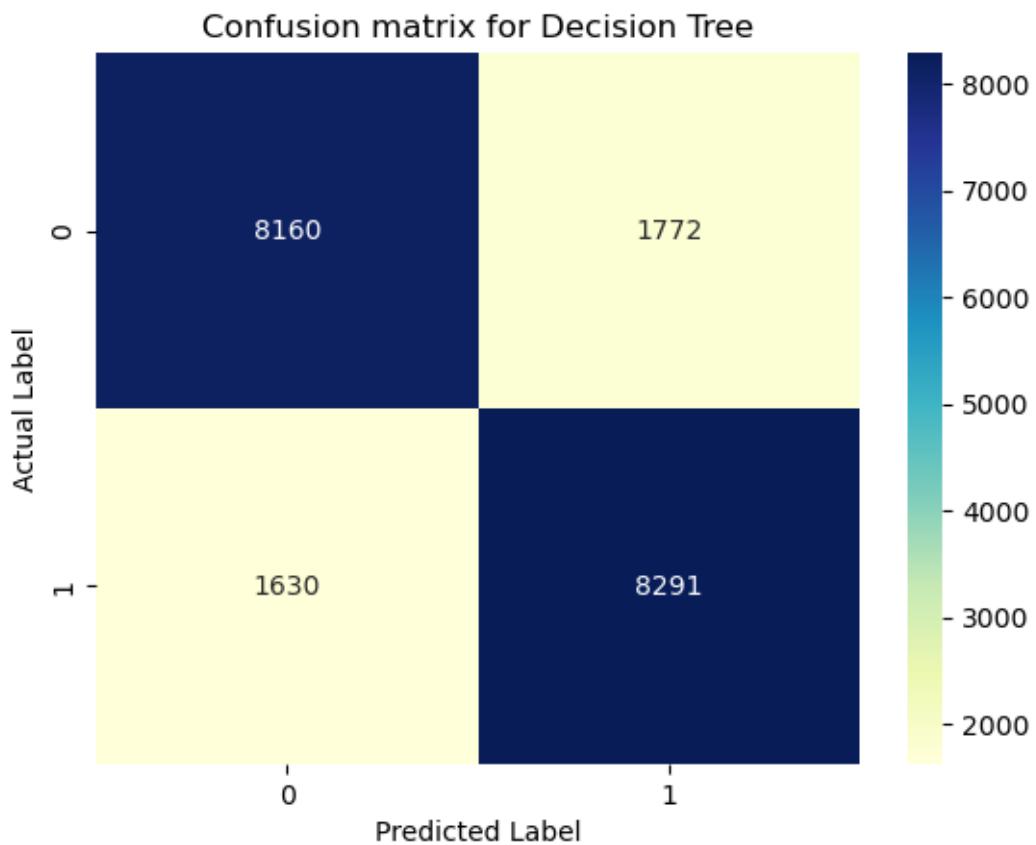


```
[51]: print ("Accuracy of Naive Bayes for training datatse is:", nb_model.
    score(X_train, y_train))
```

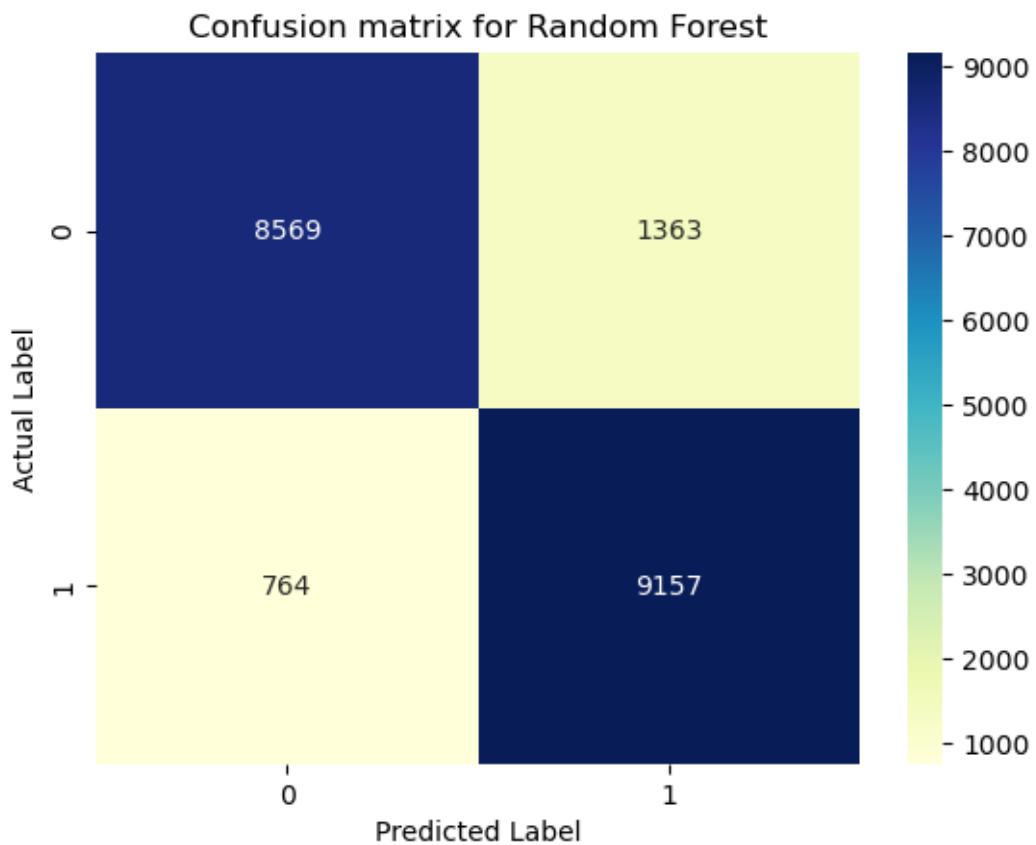
Accuracy of Naive Bayes for training datatse is: 0.7413052395241899

```
[54]: # confusion matrix for Decision Tree

sns.heatmap(cm2, annot=True, cmap='YlGnBu', fmt='g')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.title('Confusion matrix for Decision Tree');
```

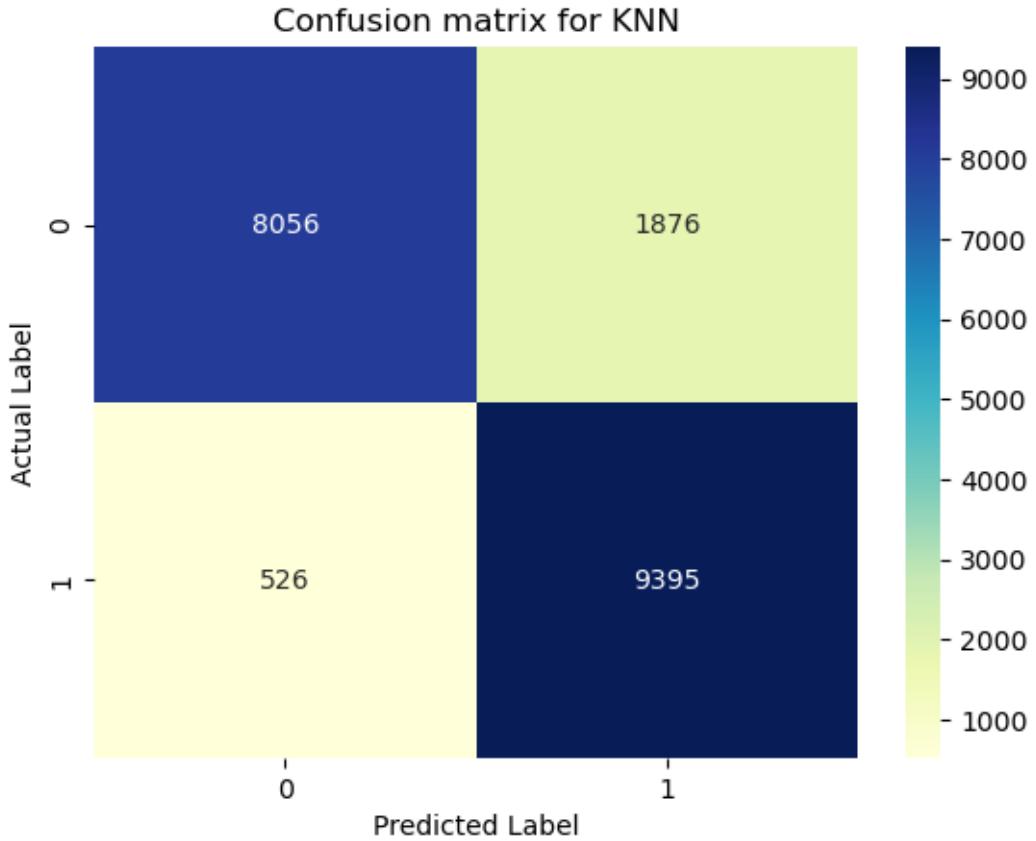


```
[55]: # confusion matrix for Random Forest  
  
sns.heatmap(cm3, annot=True, cmap='YlGnBu', fmt='g')  
plt.ylabel('Actual Label')  
plt.xlabel('Predicted Label')  
plt.title('Confusion matrix for Random Forest');
```



```
[56]: # confusion matrix for KNN

sns.heatmap(cm4, annot=True, cmap='YlGnBu', fmt='g')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.title('Confusion matrix for KNN');
```



```
[57]: # Random forest has the best performance. let's all the measures in a report
print(classification_report(y_test, y_pred_rf))
```

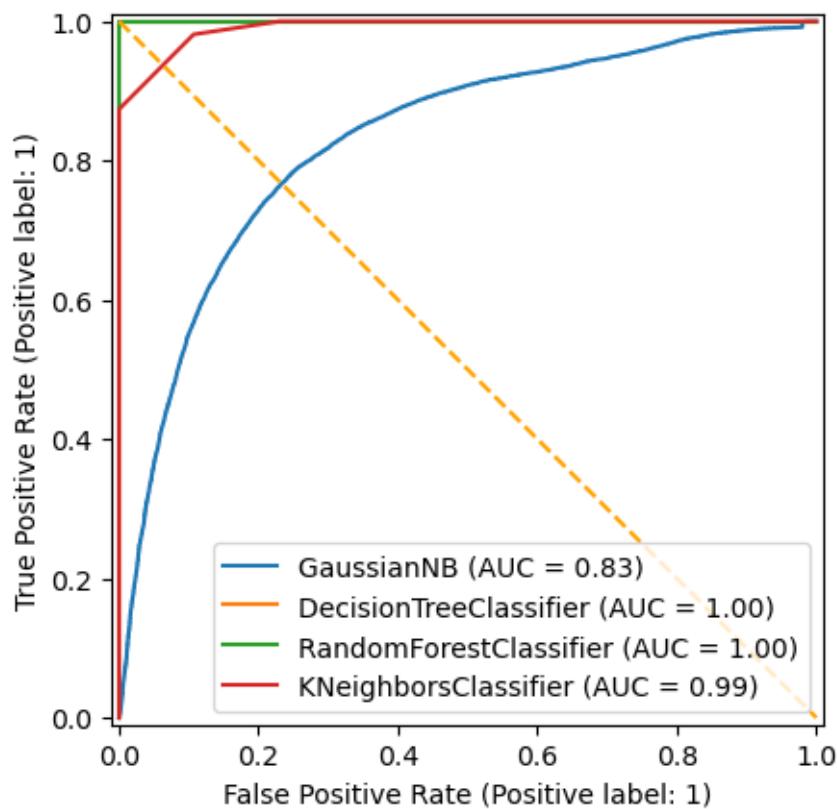
	precision	recall	f1-score	support
0	0.92	0.86	0.89	9932
1	0.87	0.92	0.90	9921
accuracy			0.89	19853
macro avg	0.89	0.89	0.89	19853
weighted avg	0.89	0.89	0.89	19853

```
[58]: # plot ROC curve for Xtrain and Xtest
```

```
from sklearn.metrics import RocCurveDisplay

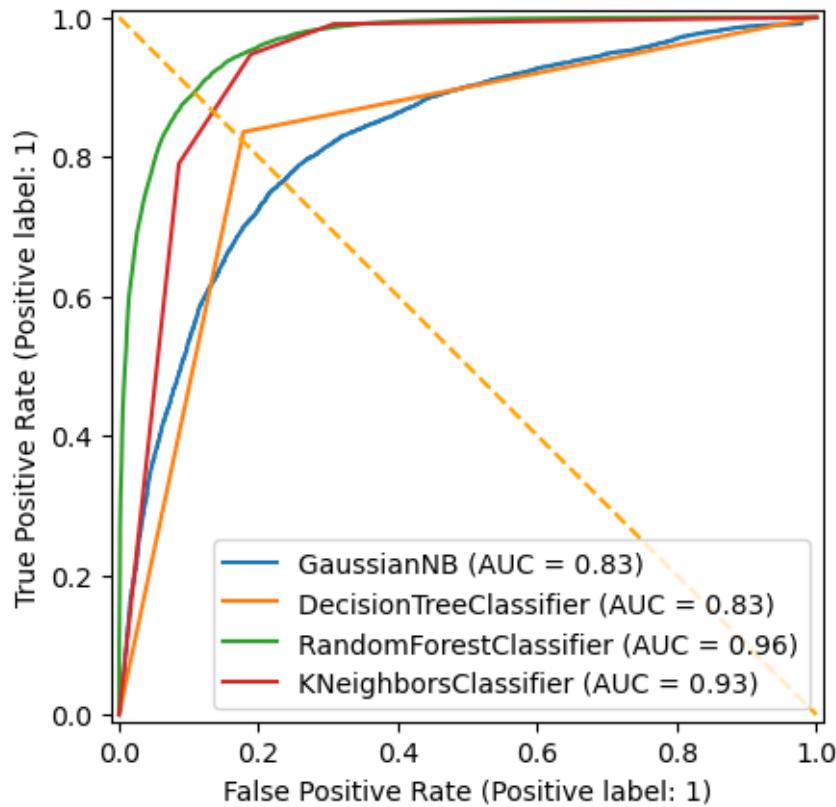
disp =RocCurveDisplay.from_estimator(nb_model, X_train, y_train)
RocCurveDisplay.from_estimator(dt_model, X_train, y_train, ax=disp.ax_)
RocCurveDisplay.from_estimator(rf_model, X_train, y_train, ax=disp.ax_)
```

```
RocCurveDisplay.from_estimator(knn_model, X_train, y_train, ax=disp.ax_)
plt.plot([0,1], [1,0], color='orange', linestyle='--');
```



[59]: # ROC Curve for Testing dataset

```
disp =RocCurveDisplay.from_estimator(nb_model, X_test, y_test)
RocCurveDisplay.from_estimator(dt_model, X_test, y_test, ax=disp.ax_)
RocCurveDisplay.from_estimator(rf_model, X_test, y_test, ax=disp.ax_)
RocCurveDisplay.from_estimator(knn_model, X_test, y_test, ax=disp.ax_)
plt.plot([0,1], [1,0], color='orange', linestyle='--');
```



```
[60]: # Make a new prediction for new clients (new observations)

# load new observations

# Define the data
data = {
    'age': [35, 28, 52, 43],
    'job': ['teacher', 'artist', 'entrepreneur', 'engineer'],
    'marital': ['single', 'divorced', 'married', 'married'],
    'education': ['tertiary', 'secondary', 'tertiary', 'secondary'],
    'default': ['no', 'no', 'yes', 'no'],
    'balance': [1200, 850, 2100, 3500],
    'housing': ['yes', 'no', 'yes', 'no'],
    'loan': ['no', 'yes', 'yes', 'no'],
    'contact': ['cellular', 'cellular', 'unknown', 'telephone'],
    'day': [12, 24, 17, 5],
    'month': ['jul', 'sep', 'aug', 'jun'],
    'duration': [320, 145, 180, 400],
    'campaign': [3, 2, 1, 4],
    'pdays': [45, -1, 60, 30],
```

```

    'previous': [1, 0, 2, 1],
    'poutcome': ['success', 'unknown', 'success', 'failure'],
    'y': ['yes', 'no', 'yes', 'no']
}

# Create the DataFrame
df_sample = pd.DataFrame(data)

# Display the DataFrame
df_sample

```

```
[60]:   age         job  marital  education default  balance housing loan \
0   35      teacher  single  tertiary     no    1200     yes    no
1   28      artist  divorced secondary     no     850     no    yes
2   52 entrepreneur married tertiary    yes    2100     yes    yes
3   43   engineer  married secondary     no    3500     no    no

      contact  day month duration campaign  pdays previous poutcome     y
0  cellular   12  jul       320        3     45          1  success  yes
1  cellular   24  sep       145        2     -1          0  unknown  no
2  unknown    17  aug       180        1      60          2  success  yes
3 telephone    5  jun       400        4     30          1 failure  no
```

```
[61]: df_sample
```

```
[61]:   age         job  marital  education default  balance housing loan \
0   35      teacher  single  tertiary     no    1200     yes    no
1   28      artist  divorced secondary     no     850     no    yes
2   52 entrepreneur married tertiary    yes    2100     yes    yes
3   43   engineer  married secondary     no    3500     no    no

      contact  day month duration campaign  pdays previous poutcome     y
0  cellular   12  jul       320        3     45          1  success  yes
1  cellular   24  sep       145        2     -1          0  unknown  no
2  unknown    17  aug       180        1      60          2  success  yes
3 telephone    5  jun       400        4     30          1 failure  no
```

```
[62]: df_sample.drop(columns='contact', inplace=True)
df_sample.drop(columns='previous', inplace=True)
df_sample.drop(columns='poutcome', inplace=True)
df_sample.drop(columns='pdays', inplace=True)
df_sample.drop(columns='month', inplace=True)
df_sample.drop(columns='day', inplace=True)
```

```
[63]: df_sample
```

```
[63]:   age          job  marital  education default  balance housing loan \
0    35      teacher    single   tertiary     no    1200     yes    no
1    28      artist  divorced secondary     no     850     no   yes
2    52  entrepreneur married tertiary    yes   2100     yes   yes
3    43     engineer married secondary     no   3500     no    no

      duration  campaign    y
0        320         3  yes
1        145         2  no
2        180         1 yes
3        400         4  no
```

1

```
[64]: lb=LabelEncoder()

for col in categorical_cols:
    df_sample.loc[:,col]=lb.fit_transform(df_sample[col])
scaler = MinMaxScaler()

df_sample.loc[:,numerical_cols] = scaler.
    ↪fit_transform(df_sample[numerical_cols])
```

```
[65]: df_sample
```

```
[65]:   age  job marital education default  balance housing loan  duration \
0  0.291667  3      2       1       0  0.132075      1     0  0.686275
1  0.000000  0      0       0       0  0.000000      0     1  0.000000
2  1.000000  2      1       1       1  0.471698      1     1  0.137255
3  0.625000  1      1       0       0  1.000000      0     0  1.000000

      campaign  y
0  0.666667  1
1  0.333333  0
2  0.000000  1
3  1.000000  0
```

```
[66]: df_sample.drop(columns='y',inplace = True)
```

```
[67]: df_sample
```

```
[67]:   age  job marital education default  balance housing loan  duration \
0  0.291667  3      2       1       0  0.132075      1     0  0.686275
1  0.000000  0      0       0       0  0.000000      0     1  0.000000
2  1.000000  2      1       1       1  0.471698      1     1  0.137255
3  0.625000  1      1       0       0  1.000000      0     0  1.000000
```

```
campaign
0 0.666667
1 0.333333
2 0.000000
3 1.000000
```

```
[68]: y_pred_new_observation_rf=rf_model.predict(df_sample)
```

```
[69]: y_pred_new_observation_nb = nb_model.predict(df_sample)
```

```
[70]: y_pred_new_observation_dt = dt_model.predict(df_sample)
```

```
[71]: y_pred_new_observation_knn = knn_model.predict(df_sample)
```

```
[72]: y_pred_new_observation_dt
```

```
[72]: array([0, 0, 0, 0])
```

```
[73]: y_pred_new_observation_knn
```

```
[73]: array([0, 0, 0, 1])
```

```
[74]: y_pred_new_observation_nb
```

```
[74]: array([1, 0, 0, 1])
```

```
[75]: y_pred_new_observation_rf
```

```
[75]: array([0, 0, 1, 1])
```

1.0.1 Feature Importance

```
[76]: # Find the importance of features ----- Features importance using Random Forest
```

```
rf=RandomForestClassifier(random_state=1, n_estimators=100)
```

```
[77]: y
```

```
0      0
1      0
2      0
3      0
4      0
..
45203  1
45205  1
45206  1
45209  0
```

```

45210      0
Name: y, Length: 37303, dtype: object

[78]: rf.fit(X,y.astype(int))

[78]: RandomForestClassifier(random_state=1)

[79]: I = rf.feature_importances_

[80]: I

[80]: array([0.16843324, 0.07493019, 0.02846266, 0.0357738 , 0.00286004,
       0.21485109, 0.02604091, 0.01490493, 0.38394217, 0.04980096])

[81]: X.columns

[81]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
       'loan', 'duration', 'campaign'],
       dtype='object')

[82]: # Define a variable that contains only the feature column names
name_features = X.columns

[83]: name_features

[83]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
       'loan', 'duration', 'campaign'],
       dtype='object')

[84]: len(name_features)

[84]: 10

[85]: len(I)

[85]: 10

[86]: # generate a new dataframe call it df_I
df_I=pd.DataFrame([I, name_features])

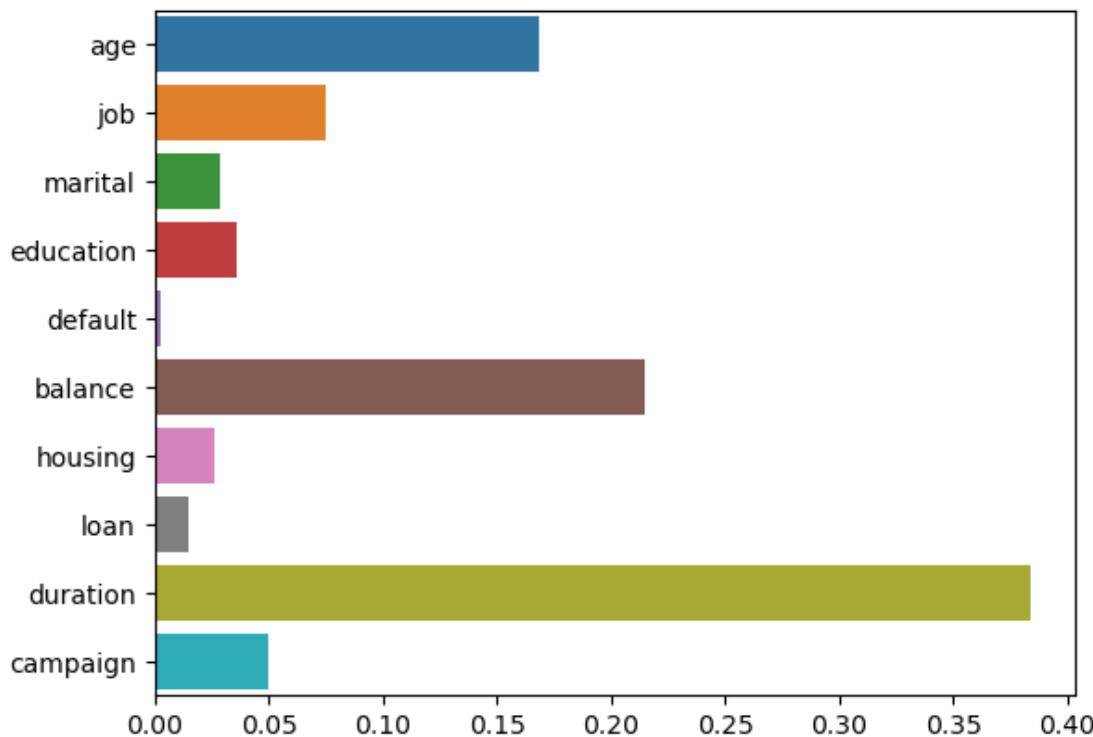
[87]: df_I

```

	0	1	2	3	4	5	6	\
0	0.168433	0.07493	0.028463	0.035774	0.00286	0.214851	0.026041	
1	age	job	marital	education	default	balance	housing	
	7	8	9					

```
0  0.014905  0.383942  0.049801  
1      loan    duration   campaign
```

```
[88]: sns.barplot(data=df_I, y=name_features, x=I);
```



1.0.2 optimal value for KNN

```
[89]: # find the best value of K in KNN  
  
from sklearn.metrics import balanced_accuracy_score
```

```
[90]: error=[]  
  
#calculating the error for k values between 1 and 20  
  
for i in range(1,21): # i=1,2,...,20  
    knn = KNeighborsClassifier(n_neighbors=i)  
    knn.fit(X_train, y_train)  
    pred_i = knn.predict(X_test)  
    accuracy = balanced_accuracy_score(y_test, pred_i)  
    error.append(1-accuracy)
```

```

error;

[91]: import matplotlib.pyplot as plt

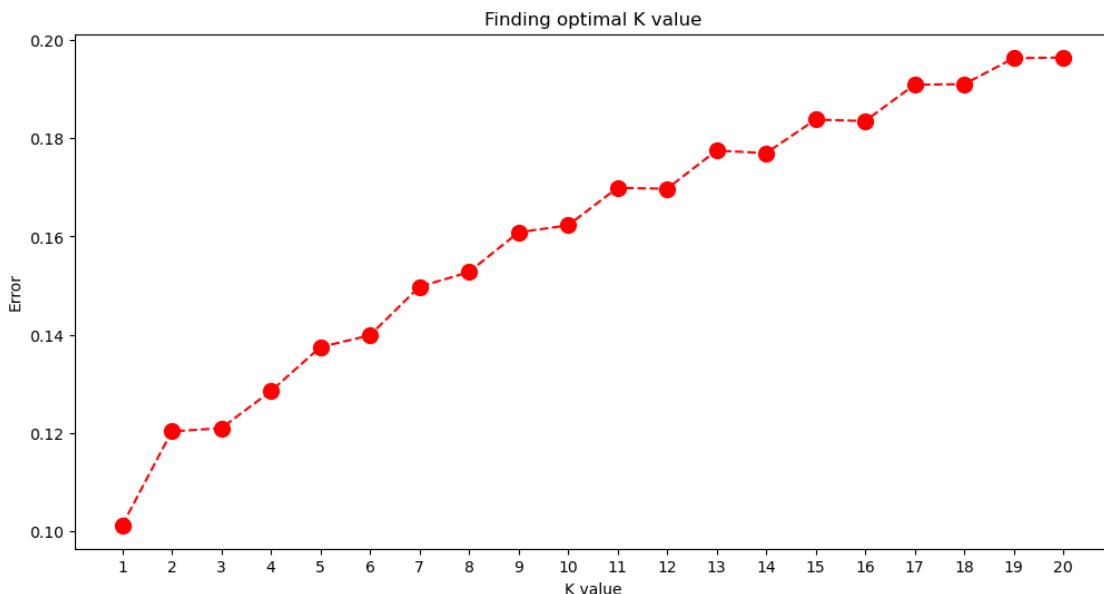
plt.figure(figsize=(12,6))

plt.plot(range(1,21), error, color='red', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)

plt.title ('Finding optimal K value')
plt.xlabel ('K value')
plt.ylabel('Error');

# Set the ticks on the X-axis to integer values
plt.xticks(ticks=range(1, 21)); # Set the ticks to integers 1 to 20

```



```
[ ]: ##### Conditional Probability #####

```

```
[4]: df_cond = pd.read_csv('bank-full.csv', delimiter=';')
```

```
[5]: df_cond
```

```
[5]:      age        job marital education default  balance housing loan \
0    58.0  management   married   tertiary     no     2143    yes   no
1    44.0  technician    single       NaN     no      29    yes   no
2    33.0 entrepreneur   married  secondary     no       2    yes  yes
3    47.0 blue-collar   married  unknown     no     1506    yes   no
```

```

4      33.0      unknown    single     unknown      no       1      no      no
...
45206  51.0  technician   married   tertiary      no      825      no      no
45207  71.0    retired   divorced   primary      no     1729      no      no
45208  72.0    retired   married   secondary      no      5715      no      no
45209  57.0  blue-collar married   secondary      no      668      no      no
45210  37.0 entrepreneur married   secondary      no     2971      no      no

      contact  day month duration campaign  pdays previous poutcome      y
0      unknown    5   may      261         1      -1          0  unknown    no
1      unknown    5   may      151         1      -1          0  unknown    no
2      unknown    5   may       76         1      -1          0  unknown    no
3      unknown    5   may       92         1      -1          0  unknown    no
4      unknown    5   may      198         1      -1          0  unknown    no
...
45206  cellular   17  nov      977         3      -1          0  unknown   yes
45207  cellular   17  nov      456         2      -1          0  unknown   yes
45208  cellular   17  nov     1127         5     184          3  success   yes
45209  telephone   17  nov      508         4      -1          0  unknown    no
45210  cellular   17  nov      361         2     188         11  other     no

```

[45211 rows x 17 columns]

```
[8]: ((df_cond['loan']=='yes') & (df_cond['y']=='yes')).sum()
```

```
[8]: 484
```

```
[10]: # the probability that a person takes term deposit given he has loan???
prob_loan_yes = (df_cond[df_cond["loan"]=="yes"]["y"] == "yes").sum() /
    ↪(df_cond[df_cond["loan"] == 'yes']["loan"]).count()
prob_loan_yes.round(3)
```

```
[10]: 0.067
```

```
[68]: # the probability that a person makes a term deposit given he has taken housing
    ↪loan ???
prob_loan_yes_housing = (df_cond[df_cond["housing"]=="yes"]["y"] == "yes").
    ↪sum() / (df_cond[df_cond["housing"] == 'yes']["housing"]).count()
prob_loan_yes_housing.round(3)
```

```
[68]: 0.077
```

```
[69]: # the probability of a person making term deposit given he has taken both
    ↪housing loan as well as normal loan???
```

```

prob_both_yes = (df_cond[(df_cond["housing"] == "yes") & (df_cond["loan"] == "yes")]
                  ["y"].sum() / len(df_cond[(df_cond["housing"] == "yes") &
                  (df_cond["loan"] == "yes")]))  

prob_both_yes.round(3)

```

[69]: 0.061

[70]: # the probability that he makes term deposit given his age is less than 40??
prob_loan_yes_age = (df_cond[df_cond["age"]<40]["y"].sum() /
(df_cond[df_cond["age"] < 40]["age"]).count())
prob_loan_yes_age.round(3)

[70]: 0.122

[71]: df_cond

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...	
45206	51	technician	married	tertiary	no	825	no	no	
45207	71	retired	divorced	primary	no	1729	no	no	
45208	72	retired	married	secondary	no	5715	no	no	
45209	57	blue-collar	married	secondary	no	668	no	no	
45210	37	entrepreneur	married	secondary	no	2971	no	no	
	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no
...	
45206	cellular	17	nov	977	3	-1	0	unknown	yes
45207	cellular	17	nov	456	2	-1	0	unknown	yes
45208	cellular	17	nov	1127	5	184	3	success	yes
45209	telephone	17	nov	508	4	-1	0	unknown	no
45210	cellular	17	nov	361	2	188	11	other	no

[45211 rows x 17 columns]

[72]: # the probability that a person has a housing loan given that he/she is single

```

prob_housing_loan_yes_age = (df_cond[df_cond["marital"]=="single"]["housing"][
    ↪== "yes").sum()/(df_cond[df_cond["marital"] == "single"]["marital"]).count()
prob_housing_loan_yes_age.round(3)

```

[72]: 0.54

```

[15]: # The feature importance of loan
FI_loan=(df_cond[df_cond["loan"]=="yes"]["y"] == "yes").sum()/
    ↪(df_cond["loan"]== "yes").sum() - (df_cond[df_cond["loan"]=="no"]["y"] ==
    ↪"yes").sum()/(df_cond["loan"]== "no").sum()
FI_loan

```

[15]: -0.05974335845754743

```

[16]: # The feature importance of housing
FI_housing=(df_cond[df_cond["housing"]=="yes"]["y"] == "yes").sum()/
    ↪(df_cond["housing"]== "yes").sum() - (df_cond[df_cond["housing"]=="no"]["y"] ==
    ↪"yes").sum()/(df_cond["housing"]== "no").sum()
FI_housing

```

[16]: -0.09002395253461445

[73]: df_linear=pd.read_csv("House Price India.csv")

[74]: df_linear

	number of bedrooms	number of bathrooms	living area	lot area	\
0	5	2.50	3650	9050	
1	4	2.50	2920	4000	
2	5	2.75	2910	9480	
3	4	2.50	3310	42998	
4	3	2.00	2710	4500	
...	
14615	2	1.50	1556	20000	
14616	3	2.00	1680	7000	
14617	2	1.00	1070	6120	
14618	4	1.00	1030	6621	
14619	3	1.00	900	4770	

	number of floors	waterfront present	number of views	\
0	2.0	0	4	
1	1.5	0	0	
2	1.5	0	0	
3	2.0	0	0	
4	1.5	0	0	
...	
14615	1.0	0	0	

14616	1.5	0	0
14617	1.0	0	0
14618	1.0	0	0
14619	1.0	0	0

	condition of the house	grade of the house	\
0	5	10	
1	5	8	
2	3	8	
3	3	9	
4	4	8	
...	
14615	4	7	
14616	4	7	
14617	3	6	
14618	4	6	
14619	3	6	

	Area of the house(excluding basement)	Area of the basement	\
0	3370	280	
1	1910	1010	
2	2910	0	
3	3310	0	
4	1880	830	
...	
14615	1556	0	
14616	1680	0	
14617	1070	0	
14618	1030	0	
14619	900	0	

	Built Year	living_area_renov	lot_area_renov	\
0	1921	2880	5400	
1	1909	2470	4000	
2	1939	2940	6600	
3	2001	3350	42847	
4	1929	2060	4500	
...	
14615	1957	2250	17286	
14616	1968	1540	7480	
14617	1962	1130	6120	
14618	1955	1420	6631	
14619	1969	900	3480	

	Number of schools nearby	Distance from the airport	Price
0	2	58	2380000
1	2	51	1400000

2		1		53	1200000
3		3		76	838000
4		1		51	805000
...
14615		3		76	221700
14616		3		59	219200
14617		2		64	209000
14618		3		54	205000
14619		2		55	146000

[14620 rows x 17 columns]

[75]: df_linear.head(5)

	number of bedrooms	number of bathrooms	living area	lot area	\
0	5	2.50	3650	9050	
1	4	2.50	2920	4000	
2	5	2.75	2910	9480	
3	4	2.50	3310	42998	
4	3	2.00	2710	4500	
	number of floors	waterfront present	number of views	\	
0	2.0	0	4		
1	1.5	0	0		
2	1.5	0	0		
3	2.0	0	0		
4	1.5	0	0		
	condition of the house	grade of the house	\		
0	5	10			
1	5	8			
2	3	8			
3	3	9			
4	4	8			
	Area of the house(excluding basement)	Area of the basement	Built Year	\	
0		3370	280	1921	
1		1910	1010	1909	
2		2910	0	1939	
3		3310	0	2001	
4		1880	830	1929	
	living_area_renov	lot_area_renov	Number of schools nearby	\	
0	2880	5400	2		
1	2470	4000	2		
2	2940	6600	1		
3	3350	42847	3		

```
4           2060        4500          1
```

	Distance from the airport	Price
0	58	2380000
1	51	1400000
2	53	1200000
3	76	838000
4	51	805000

```
[76]: df_linear.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 17 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   number of bedrooms    14620 non-null   int64  
 1   number of bathrooms   14620 non-null   float64 
 2   living area           14620 non-null   int64  
 3   lot area              14620 non-null   int64  
 4   number of floors      14620 non-null   float64 
 5   waterfront present   14620 non-null   int64  
 6   number of views       14620 non-null   int64  
 7   condition of the house 14620 non-null   int64  
 8   grade of the house    14620 non-null   int64  
 9   Area of the house(excluding basement) 14620 non-null   int64  
 10  Area of the basement 14620 non-null   int64  
 11  Built Year           14620 non-null   int64  
 12  living_area_renov    14620 non-null   int64  
 13  lot_area_renov       14620 non-null   int64  
 14  Number of schools nearby 14620 non-null   int64  
 15  Distance from the airport 14620 non-null   int64  
 16  Price                14620 non-null   int64  
dtypes: float64(2), int64(15)
memory usage: 1.9 MB
```

```
[77]: X=df_linear.iloc[:,0:14]
y=df_linear["Price"]
X1=df_linear.iloc[:,0].values
```

```
[80]: lin_reg=LinearRegression()
```

```
[81]: lin_reg
```

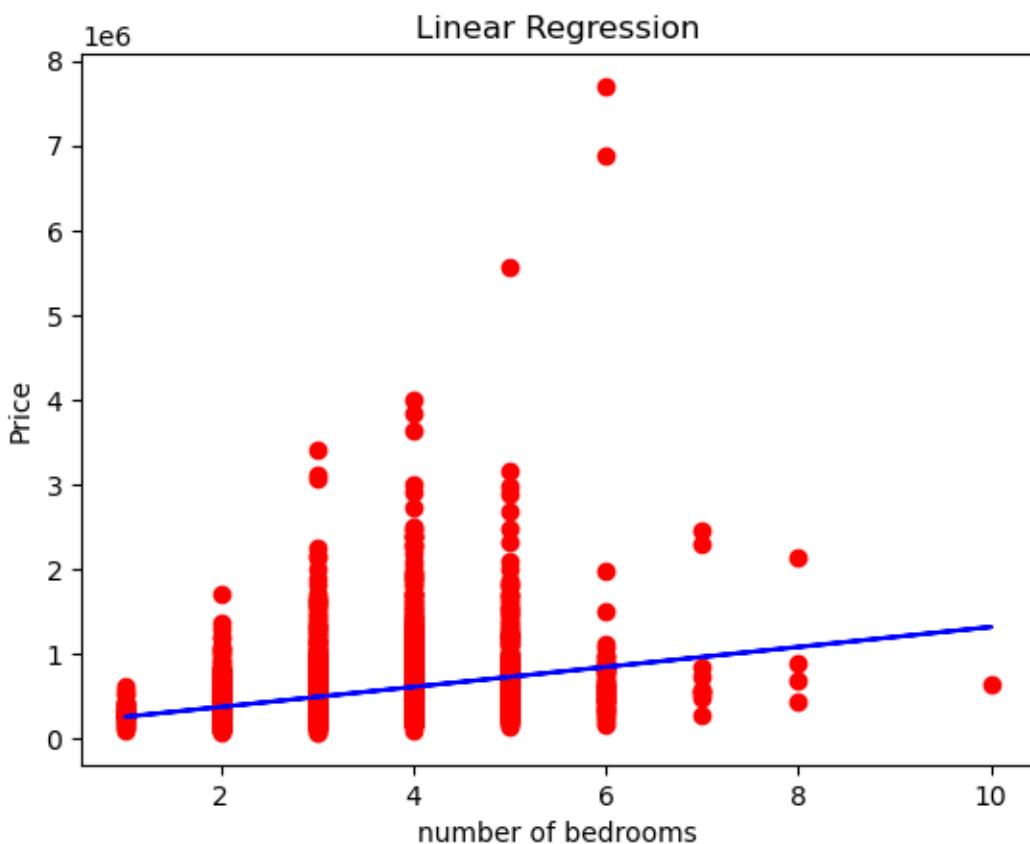
```
[81]: LinearRegression()
```

```
[82]: Xtrain, Xtest, ytrain, ytest = train_test_split (X1, y, test_size=0.3,
random_state=0)
```

```
lin_reg.fit(Xtrain.reshape(-1,1), ytrain)
```

[82]: LinearRegression()

```
[84]: ypred=lin_reg.predict(Xtest.reshape(-1,1))
      #correlation between number of bedrooms and price
      plt.scatter(Xtest, ytest, color='red')
      plt.plot(Xtest, ypred, color='blue')
      plt.title('Linear Regression')
      plt.xlabel('number of bedrooms')
      plt.ylabel('Price');
```



```
[87]: X=df_linear.iloc[:,0:14]
      y=df_linear["Price"]
      X3=df_linear.iloc[:,2].values
      type(X3)
      X3;
```

```
[88]: lin_reg=LinearRegression()  
lin reg
```

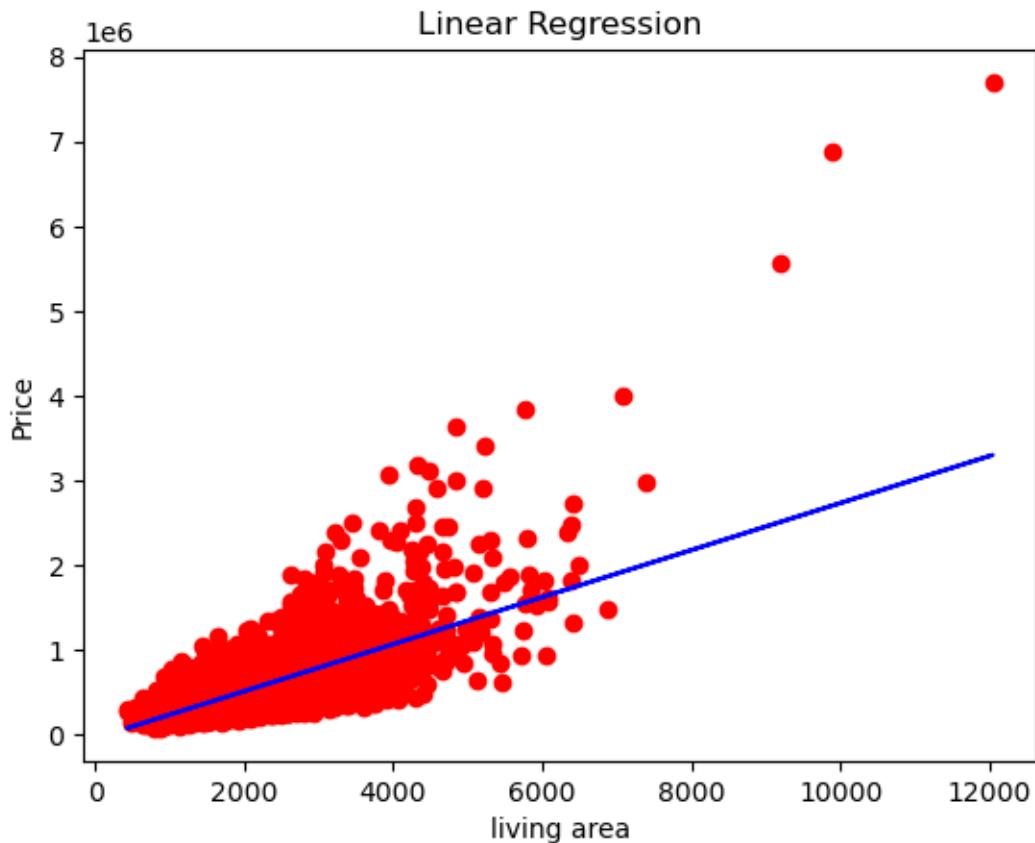
```
[88]: LinearRegression()
```

```
[89]: Xtrain, Xtest, ytrain, ytest = train_test_split (X3, y, test_size=0.3,  
random_state=0)  
lin_reg.fit(Xtrain.reshape(-1,1), ytrain)
```

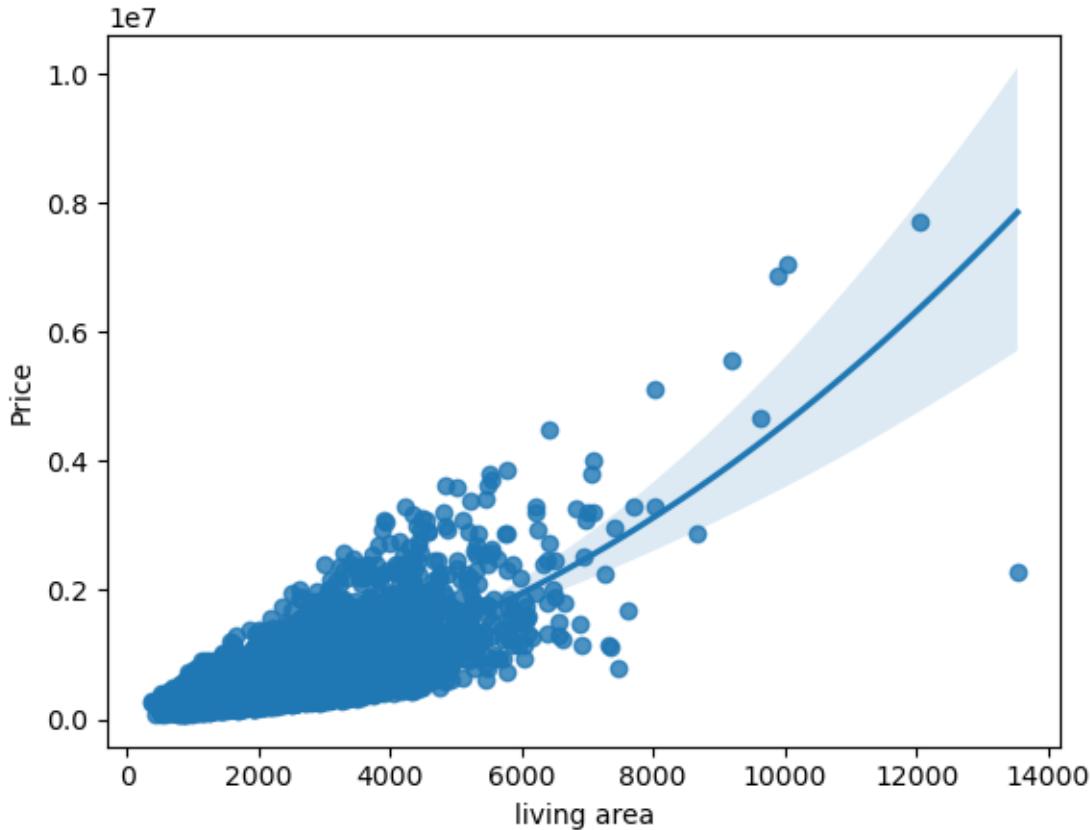
```
[89]: LinearRegression()
```

```
[90]: ypred=lin_reg.predict(Xtest.reshape(-1,1))
```

```
[91]: plt.scatter(Xtest, ytest, color='red')  
plt.plot(Xtest, ypred, color='blue')  
plt.title('Linear Regression')  
plt.xlabel('living area')  
plt.ylabel('Price');
```



```
[92]: sns.regplot(x="living area", y="Price", data=df_linear, order=2);
```



```
[94]: #correlation between grade of the house and price
X=df_linear.iloc[:,0:14]
y=df_linear["Price"]
X7=df_linear.iloc[:,6].values
type(X7)
X7;
```

```
[98]: lin_reg=LinearRegression()
lin_reg
```

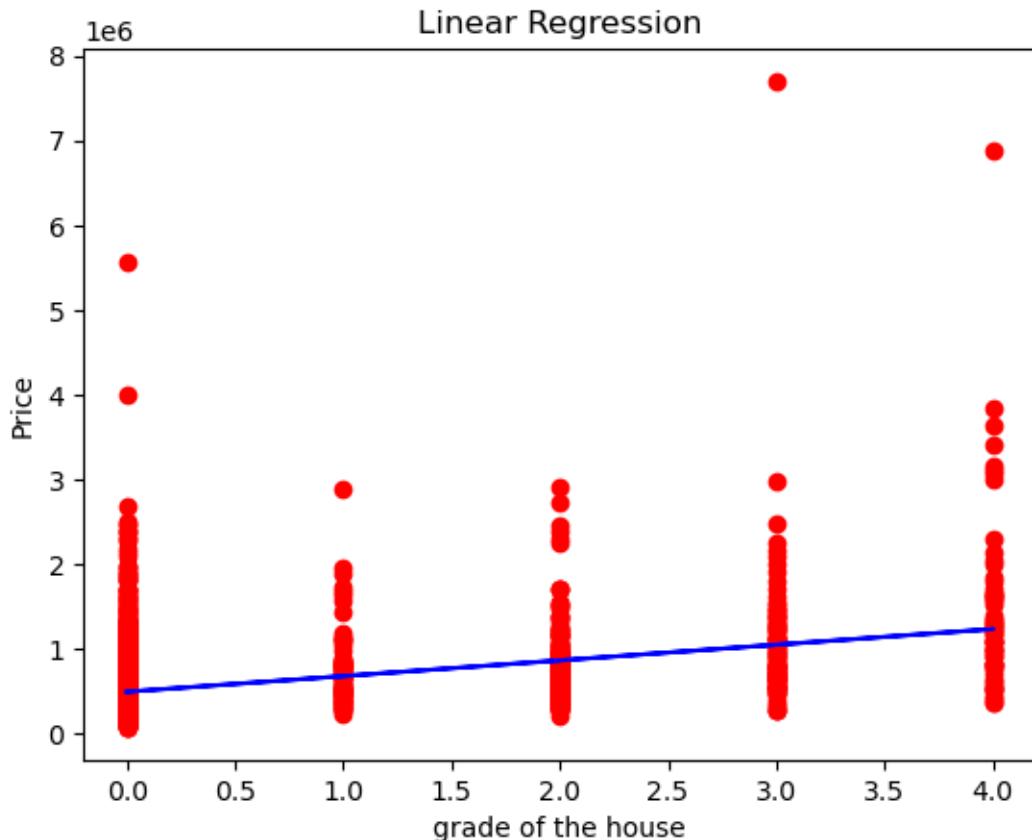
```
[98]: LinearRegression()
```

```
[99]: Xtrain, Xtest, ytrain, ytest = train_test_split (X7, y, test_size=0.3,
random_state=0)
lin_reg.fit(Xtrain.reshape(-1,1), ytrain)
```

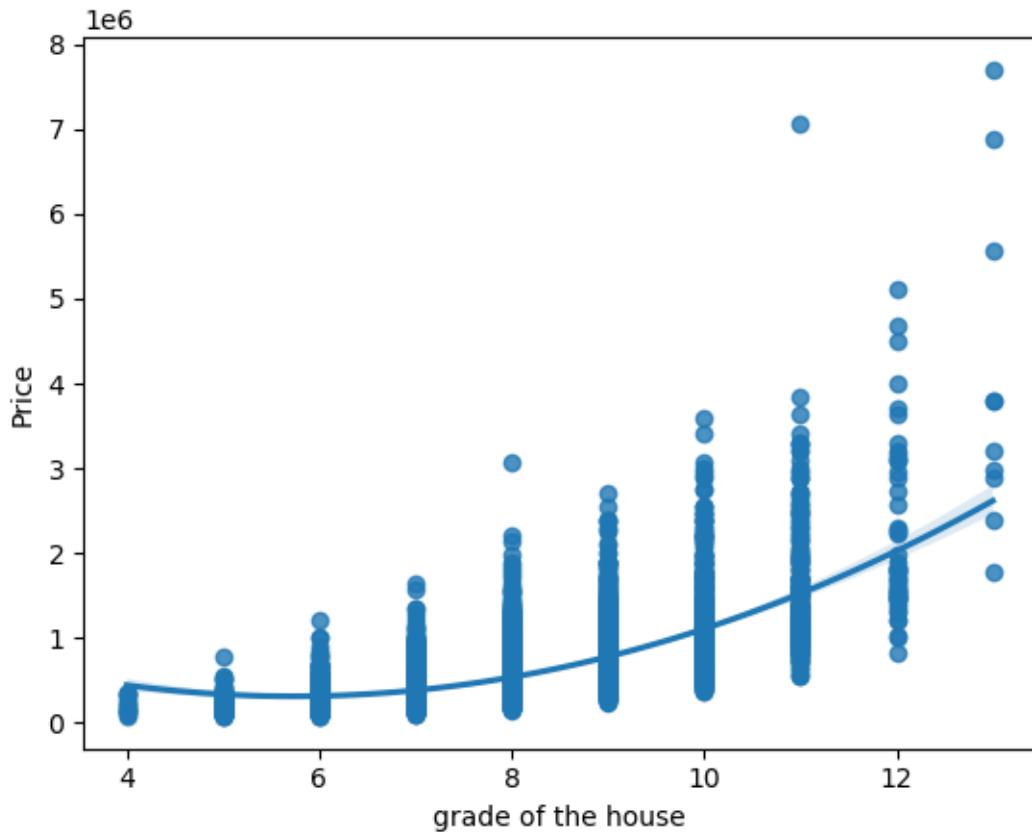
```
[99]: LinearRegression()
```

```
[100]: ypred=lin_reg.predict(Xtest.reshape(-1,1))
```

```
[101]: plt.scatter(Xtest, ytest, color='red')
plt.plot(Xtest, ypred, color='blue')
plt.title('Linear Regression')
plt.xlabel('grade of the house')
plt.ylabel('Price');
```



```
[102]: sns.regplot(x="grade of the house", y="Price", data=df_linear, order=2);
```



```
[103]: df=pd.read_csv("House Price India.csv")
```

```
[104]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   number of bedrooms    14620 non-null   int64  
 1   number of bathrooms   14620 non-null   float64 
 2   living area          14620 non-null   int64  
 3   lot area             14620 non-null   int64  
 4   number of floors     14620 non-null   float64 
 5   waterfront present   14620 non-null   int64  
 6   number of views      14620 non-null   int64  
 7   condition of the house 14620 non-null   int64  
 8   grade of the house   14620 non-null   int64  
 9   Area of the house(excluding basement) 14620 non-null   int64  
 10  Area of the basement 14620 non-null   int64  
 11  Built Year          14620 non-null   int64
```

```

12 living_area_renov           14620 non-null  int64
13 lot_area_renov             14620 non-null  int64
14 Number of schools nearby  14620 non-null  int64
15 Distance from the airport 14620 non-null  int64
16 Price                      14620 non-null  int64
dtypes: float64(2), int64(15)
memory usage: 1.9 MB

```

[105]: df.head(5)

```

[105]:   number of bedrooms  number of bathrooms  living area  lot area  \
0                  5              2.50        3650      9050
1                  4              2.50        2920      4000
2                  5              2.75        2910      9480
3                  4              2.50        3310     42998
4                  3              2.00        2710      4500

      number of floors  waterfront present  number of views  \
0                2.0                 0            4
1                1.5                 0            0
2                1.5                 0            0
3                2.0                 0            0
4                1.5                 0            0

      condition of the house  grade of the house  \
0                  5                  10
1                  5                  8
2                  3                  8
3                  3                  9
4                  4                  8

      Area of the house(excluding basement)  Area of the basement  Built Year  \
0                           3370                  280       1921
1                           1910                 1010       1909
2                           2910                  0       1939
3                           3310                  0       2001
4                           1880                 830       1929

      living_area_renov  lot_area_renov  Number of schools nearby  \
0                  2880          5400                  2
1                  2470          4000                  2
2                  2940          6600                  1
3                  3350         42847                  3
4                  2060          4500                  1

      Distance from the airport      Price
0                         58    2380000

```

```

1           51  1400000
2           53  1200000
3           76   838000
4           51   805000

```

```
[106]: # Calculate the correlation between input variables, as well as between inputs ↵and output variables.

# Calculate the correlation matrix for the dataset
correlation_matrix = df.corr()
correlation_matrix
```

	number of bedrooms \
number of bedrooms	1.000000
number of bathrooms	0.509784
living area	0.570526
lot area	0.034416
number of floors	0.177294
waterfront present	-0.006257
number of views	0.078665
condition of the house	0.026597
grade of the house	0.352945
Area of the house(excluding basement)	0.473599
Area of the basement	0.300332
Built Year	0.152954
living_area_renov	0.389855
lot_area_renov	0.029400
Number of schools nearby	0.003397
Distance from the airport	-0.006157
Price	0.308460

	number of bathrooms	living area \
number of bedrooms	0.509784	0.570526
number of bathrooms	1.000000	0.753517
living area	0.753517	1.000000
lot area	0.080806	0.174420
number of floors	0.502924	0.354743
waterfront present	0.060104	0.105837
number of views	0.183789	0.287728
condition of the house	-0.128232	-0.063358
grade of the house	0.663054	0.761835
Area of the house(excluding basement)	0.684391	0.875793
Area of the basement	0.287190	0.441491
Built Year	0.498127	0.309602
living_area_renov	0.570530	0.757571
lot_area_renov	0.078627	0.180312
Number of schools nearby	0.002180	0.002370
Distance from the airport	0.009206	0.002511

Price	0.531735	0.712169
number of bedrooms	lot area	number of floors \
number of bathrooms	0.034416	0.177294
living area	0.080806	0.502924
lot area	0.174420	0.354743
number of floors	1.000000	-0.004138
waterfront present	-0.004138	1.000000
number of views	0.026282	0.016316
condition of the house	0.078308	0.020153
grade of the house	-0.008548	-0.269928
Area of the house(excluding basement)	0.110546	0.463082
Area of the basement	0.183553	0.525643
Built Year	0.019755	-0.242976
living_area_renov	0.051615	0.481565
lot_area_renov	0.149744	0.285093
Number of schools nearby	0.706812	-0.010120
Distance from the airport	-0.012671	-0.007579
Price	0.003291	0.016567
	0.081992	0.262732
number of bedrooms	waterfront present	number of views \
number of bathrooms	-0.006257	0.078665
living area	0.060104	0.183789
lot area	0.105837	0.287728
number of floors	0.026282	0.078308
waterfront present	0.016316	0.020153
number of views	1.000000	0.400206
condition of the house	0.400206	1.000000
grade of the house	0.018644	0.052533
Area of the house(excluding basement)	0.079831	0.254532
Area of the basement	0.071865	0.162672
Built Year	0.085441	0.293062
living_area_renov	-0.024226	-0.055357
lot_area_renov	0.085743	0.281452
Number of schools nearby	0.032055	0.072300
Distance from the airport	0.001563	0.008004
Price	0.001448	-0.001657
	0.263687	0.395973
number of bedrooms	condition of the house	\
number of bathrooms	0.026597	
living area	-0.128232	
lot area	-0.063358	
number of floors	-0.008548	
waterfront present	-0.269928	
	0.018644	

number of views	0.052533
condition of the house	1.000000
grade of the house	-0.152530
Area of the house(excluding basement)	-0.167695
Area of the basement	0.180609
Built Year	-0.381718
living_area_renov	-0.099743
lot_area_renov	-0.004748
Number of schools nearby	-0.006939
Distance from the airport	-0.002136
Price	0.041376

grade of the house \

number of bedrooms	0.352945
number of bathrooms	0.663054
living area	0.761835
lot area	0.110546
number of floors	0.463082
waterfront present	0.079831
number of views	0.254532
condition of the house	-0.152530
grade of the house	1.000000
Area of the house(excluding basement)	0.758222
Area of the basement	0.167160
Built Year	0.440358
living_area_renov	0.720019
lot_area_renov	0.116725
Number of schools nearby	0.000986
Distance from the airport	0.004940
Price	0.671814

Area of the house(excluding basement) \

number of bedrooms	0.473599
number of bathrooms	0.684391
living area	0.875793
lot area	0.183553
number of floors	0.525643
waterfront present	0.071865
number of views	0.162672
condition of the house	-0.167695
grade of the house	0.758222
Area of the house(excluding basement)	1.000000
Area of the basement	-0.046445
Built Year	0.419369
living_area_renov	0.737744
lot_area_renov	0.194670
Number of schools nearby	-0.002894

Distance from the airport	0.001222
Price	0.615220

	Area of the basement	Built Year	\
number of bedrooms	0.300332	0.152954	
number of bathrooms	0.287190	0.498127	
living area	0.441491	0.309602	
lot area	0.019755	0.051615	
number of floors	-0.242976	0.481565	
waterfront present	0.085441	-0.024226	
number of views	0.293062	-0.055357	
condition of the house	0.180609	-0.381718	
grade of the house	0.167160	0.440358	
Area of the house(excluding basement)	-0.046445	0.419369	
Area of the basement	1.000000	-0.138843	
Built Year	-0.138843	1.000000	
living_area_renov	0.196403	0.328625	
lot_area_renov	0.011283	0.072874	
Number of schools nearby	0.010284	-0.001631	
Distance from the airport	0.002926	-0.003968	
Price	0.330202	0.050307	

	living_area_renov	lot_area_renov	\
number of bedrooms	0.389855	0.029400	
number of bathrooms	0.570530	0.078627	
living area	0.757571	0.180312	
lot area	0.149744	0.706812	
number of floors	0.285093	-0.010120	
waterfront present	0.085743	0.032055	
number of views	0.281452	0.072300	
condition of the house	-0.099743	-0.004748	
grade of the house	0.720019	0.116725	
Area of the house(excluding basement)	0.737744	0.194670	
Area of the basement	0.196403	0.011283	
Built Year	0.328625	0.072874	
living_area_renov	1.000000	0.189225	
lot_area_renov	0.189225	1.000000	
Number of schools nearby	-0.001203	-0.025014	
Distance from the airport	-0.005673	-0.014587	
Price	0.584924	0.075535	

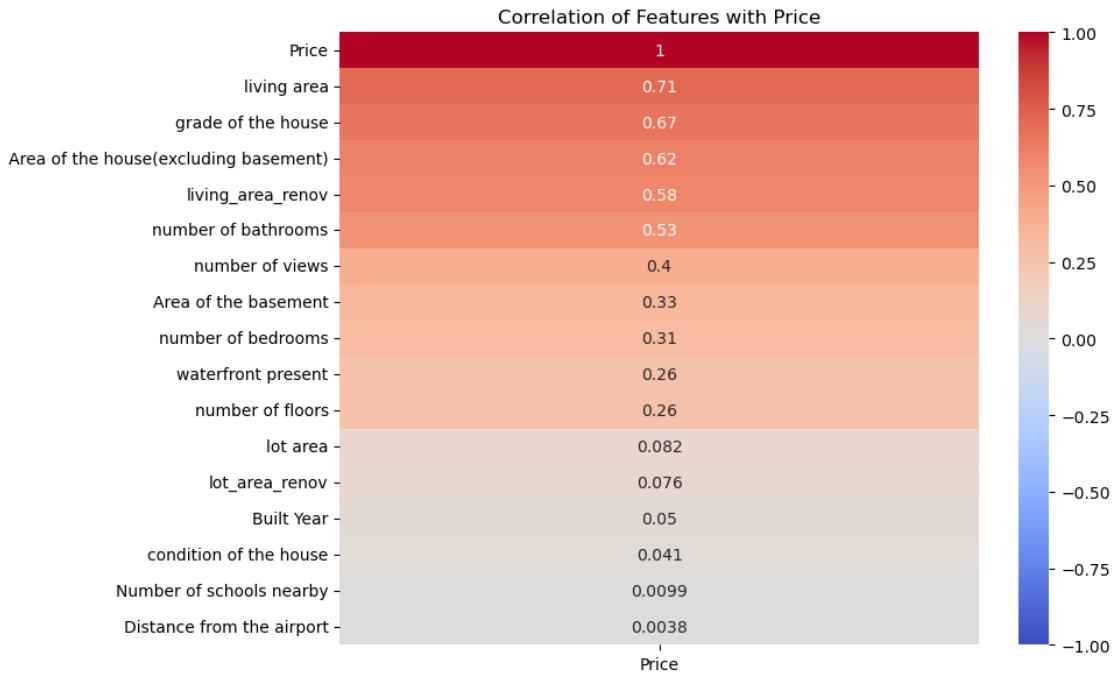
	Number of schools nearby	\
number of bedrooms	0.003397	
number of bathrooms	0.002180	
living area	0.002370	
lot area	-0.012671	
number of floors	-0.007579	

waterfront present	0.001563
number of views	0.008004
condition of the house	-0.006939
grade of the house	0.000986
Area of the house(excluding basement)	-0.002894
Area of the basement	0.010284
Built Year	-0.001631
living_area_renov	-0.001203
lot_area_renov	-0.025014
Number of schools nearby	1.000000
Distance from the airport	0.004035
Price	0.009890

	Distance from the airport	Price
number of bedrooms	-0.006157	0.308460
number of bathrooms	0.009206	0.531735
living area	0.002511	0.712169
lot area	0.003291	0.081992
number of floors	0.016567	0.262732
waterfront present	0.001448	0.263687
number of views	-0.001657	0.395973
condition of the house	-0.002136	0.041376
grade of the house	0.004940	0.671814
Area of the house(excluding basement)	0.001222	0.615220
Area of the basement	0.002926	0.330202
Built Year	-0.003968	0.050307
living_area_renov	-0.005673	0.584924
lot_area_renov	-0.014587	0.075535
Number of schools nearby	0.004035	0.009890
Distance from the airport	1.000000	0.003804
Price	0.003804	1.000000

```
[107]: plt.figure(figsize=(10, 6)) # Adjust width and height as needed
sns.heatmap(correlation_matrix[['Price']].sort_values(by='Price', ↴
    ascending=False),
            annot=True, cmap='coolwarm', vmin=-1, vmax=1, cbar=True)
plt.yticks(rotation=0)

plt.title('Correlation of Features with Price')
plt.tight_layout() # Adjusts the plot to fit within the figure area
plt.show()
```



[108]: #Living area has the highest positive correlation with price (0.71), indicating a strong relationship.
#Grade of the house also has a significant positive correlation with price (0.67).
#Area of the house (excluding basement) and living_area_renov show moderate positive correlations with price (0.61 and 0.58, respectively).
#Features like number of bathrooms (0.53), area of the basement (0.33), and number of bedrooms (0.31) have lower but still relevant positive correlations with price.
#Features like lot area, built year, condition of the house, and distance from the airport have minimal correlation with price.

[109]: X = df.drop(columns=["Price"])
y = df['Price']

[110]: # Splitting the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
Initialize and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

[110]: LinearRegression()

```
[111]: # Make predictions on the test set
y_pred = model.predict(X_test)
coefficients = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
intercept = model.intercept_

[112]: mae = mean_absolute_error(y_test, y_pred)

[115]: X = sm.add_constant(X)
reg = sm.OLS(y, X).fit()
print(reg.summary())
# output
```

OLS Regression Results

Dep. Variable:	Price	R-squared:	0.664	
Model:	OLS	Adj. R-squared:	0.664	
Method:	Least Squares	F-statistic:	1926.	
Date:	Fri, 04 Oct 2024	Prob (F-statistic):	0.00	
Time:	23:46:26	Log-Likelihood:	-2.0012e+05	
No. Observations:	14620	AIC:	4.003e+05	
Df Residuals:	14604	BIC:	4.004e+05	
Df Model:	15			
Covariance Type:	nonrobust			
P> t	[0.025 0.975]	coef	std err	t
const		6.156e+06	1.57e+05	39.153
0.000	5.85e+06	6.46e+06		
number of bedrooms		-3.984e+04	2387.221	-16.688
0.000	-4.45e+04	-3.52e+04		
number of bathrooms		4.652e+04	4141.671	11.232
0.000	3.84e+04	5.46e+04		
living area		114.4892	2.876	39.814
0.000	108.853	120.126		
lot area		-0.1311	0.066	-1.985
0.047	-0.261	-0.002		
number of floors		2.474e+04	4536.891	5.454
0.000	1.59e+04	3.36e+04		
waterfront present		5.77e+05	2.21e+04	26.078
0.000	5.34e+05	6.2e+05		
number of views		4.217e+04	2736.267	15.410
0.000	3.68e+04	4.75e+04		
condition of the house		2.151e+04	2921.274	7.364
0.000	1.58e+04	2.72e+04		
grade of the house		1.189e+05	2712.828	43.834

0.000	1.14e+05	1.24e+05			
Area of the house(excluding basement)			58.5015	2.827	20.691
0.000	52.959	64.043			
Area of the basement			55.9878	3.313	16.901
0.000	49.495	62.481			
Built Year			-3550.9701	80.317	-44.212
0.000	-3708.402	-3393.538			
living_area_renov			13.0214	4.328	3.009
0.003	4.538	21.504			
lot_area_renov			-0.5166	0.097	-5.331
0.000	-0.707	-0.327			
Number of schools nearby			3040.0393	2157.430	1.409
0.159	-1188.796	7268.875			
Distance from the airport			-113.8137	197.351	-0.577
0.564	-500.647	273.019			
<hr/>					
Omnibus:	11983.201	Durbin-Watson:		1.624	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		1201714.919	
Skew:	3.347	Prob(JB):		0.00	
Kurtosis:	46.908	Cond. No.		2.95e+17	
<hr/>					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 3.76e-22. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[116]: df['estimated']=reg.predict(X)
y=df['Price']
yhat=df['estimated']

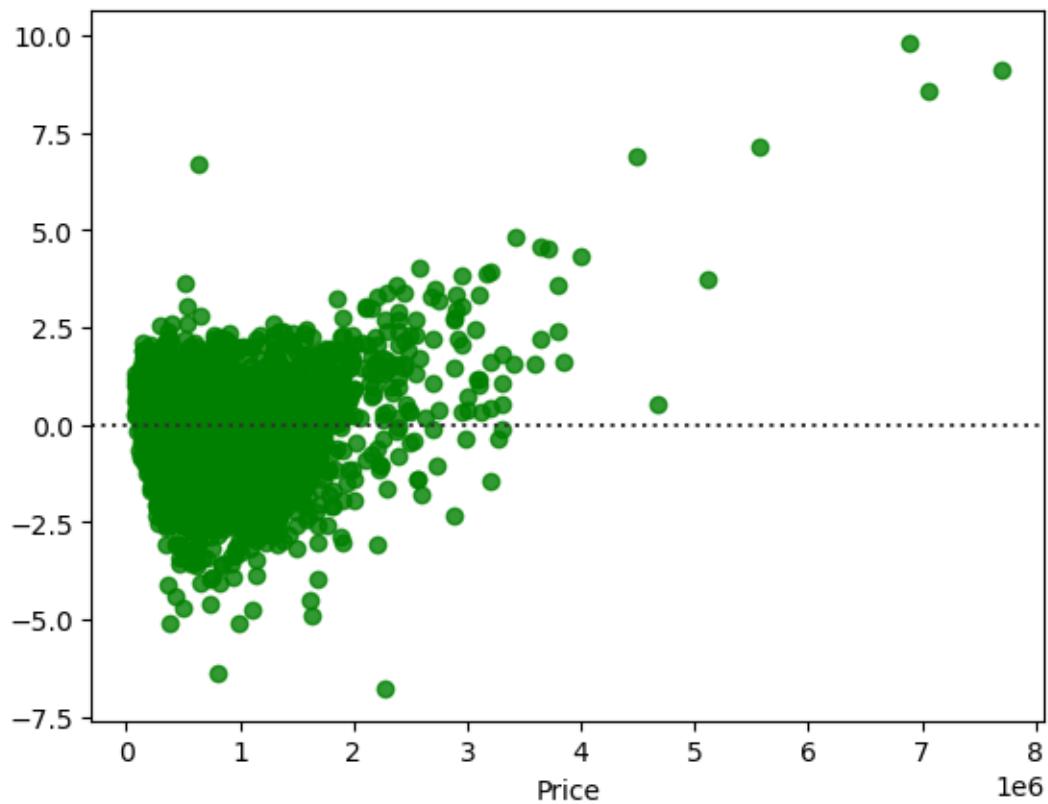
residual = y-yhat

df['residual']=residual
```

```
[119]: res= stat()
res.reg_metric(y=np.array(y), yhat=np.array(reg.predict(X)), resid=np.array(reg.
    ↴resid))
res.reg_metric_df
```

	Metrics	Value
0	Root Mean Square Error (RMSE)	2.129635e+05
1	Mean Squared Error (MSE)	4.535346e+10
2	Mean Absolute Error (MAE)	1.366337e+05
3	Mean Absolute Percentage Error (MAPE)	2.851000e-01

```
[120]: influence=reg.get_influence()
residual_price=influence.resid_studentized_external
sns.residplot(x = 'Price',y = residual_price,data = df,color='g');
```



1.0.3 Polynomial Regression

```
[123]: poly = PolynomialFeatures(degree=2)
X_poly_train = poly.fit_transform(X_train)
X_poly_test = poly.transform(X_test)
```

```
[124]: poly_model = LinearRegression()
poly_model.fit(X_poly_train, y_train)
```

```
[124]: LinearRegression()
```

```
[125]: y_pred_poly = poly_model.predict(X_poly_test)
```

```
[126]: Ypred_poly = poly_model.predict(X_poly_train)
residual_poly = np.array(y_train-Ypred_poly)
```

```
[127]: res= stat()
res.reg_metric(np.array(y_train), np.array(Ypred_poly) , resid=residual_poly)
res.reg_metric_df
```

	Metrics	Value
0	Root Mean Square Error (RMSE)	1.773649e+05
1	Mean Squared Error (MSE)	3.145829e+10
2	Mean Absolute Error (MAE)	1.193315e+05
3	Mean Absolute Percentage Error (MAPE)	2.510000e-01

```
[128]: reg = sm.OLS(y_train, X_poly_train).fit()
reg.summary()
```

Dep. Variable:	Price	R-squared:	0.765
Model:	OLS	Adj. R-squared:	0.763
Method:	Least Squares	F-statistic:	283.7
Date:	Fri, 04 Oct 2024	Prob (F-statistic):	0.00
Time:	23:54:24	Log-Likelihood:	-1.5795e+05
No. Observations:	11696	AIC:	3.162e+05
Df Residuals:	11562	BIC:	3.172e+05
Df Model:	133		
Covariance Type:	nonrobust		

	coef	std err	t	P> t 	[0.025	0.975]
const	6.587e+07	1.12e+07	5.892	0.000	4.4e+07	8.78e+07
x1	7.864e+05	1.9e+05	4.141	0.000	4.14e+05	1.16e+06
x2	-1.07e+06	3.31e+05	-3.235	0.001	-1.72e+06	-4.22e+05
x3	2766.1871	845.359	3.272	0.001	1109.141	4423.234
x4	-17.0665	7.663	-2.227	0.026	-32.087	-2.047
x5	-2.047e+06	4.43e+05	-4.620	0.000	-2.92e+06	-1.18e+06
x6	-3.806e+06	1.06e+06	-3.574	0.000	-5.89e+06	-1.72e+06
x7	2.246e+05	2.23e+05	1.006	0.315	-2.13e+05	6.62e+05
x8	3.111e+05	2.51e+05	1.241	0.215	-1.8e+05	8.02e+05
x9	8.794e+05	2.22e+05	3.969	0.000	4.45e+05	1.31e+06
x10	-2410.6967	827.769	-2.912	0.004	-4033.263	-788.130
x11	-3485.0679	859.853	-4.053	0.000	-5170.525	-1799.611
x12	-7.023e+04	1.13e+04	-6.215	0.000	-9.24e+04	-4.81e+04
x13	3921.0954	356.171	11.009	0.000	3222.940	4619.251
x14	-49.7338	11.427	-4.352	0.000	-72.133	-27.335
x15	-3.435e+04	1.83e+05	-0.188	0.851	-3.92e+05	3.24e+05
x16	-1.17e+04	1.67e+04	-0.701	0.483	-4.44e+04	2.1e+04
x17	1364.5574	331.531	4.116	0.000	714.701	2014.414
x18	2632.8968	4537.348	0.580	0.562	-6261.072	1.15e+04
x19	-15.8761	3.199	-4.963	0.000	-22.147	-9.605
x20	0.1575	0.091	1.726	0.084	-0.021	0.336
x21	1.45e+04	6456.183	2.246	0.025	1843.047	2.72e+04
x22	-899.5737	2.72e+04	-0.033	0.974	-5.42e+04	5.24e+04
x23	-2490.1847	3355.248	-0.742	0.458	-9067.039	4086.670
x24	-2996.9839	3773.366	-0.794	0.427	-1.04e+04	4399.452
x25	-4022.7120	3434.962	-1.171	0.242	-1.08e+04	2710.395
x26	-1.5664	3.661	-0.428	0.669	-8.743	5.610
x27	-9.2302	4.025	-2.293	0.022	-17.121	-1.340
x28	-428.9536	98.240	-4.366	0.000	-621.521	-236.386
x29	27.2761	5.497	4.962	0.000	16.501	38.051
x30	-0.0486	0.145	-0.336	0.737	-0.332	0.235
x31	4431.9250	2936.635	1.509	0.131	-1324.377	1.02e+04
x32	383.8766	269.496	1.424	0.154	-144.380	912.134
x33	-9059.6331	5173.316	-1.751	0.080	-1.92e+04	1080.942
x34	20.8859	5.242	3.984	0.000	10.611	31.161
x35	-0.3223	0.161	-2.003	0.045	-0.638	-0.007
x36	-2.975e+04	9501.624	-3.131	0.002	-4.84e+04	-1.11e+04
x37	-6.298e+04	3.86e+04	-1.634	0.102	-1.39e+05	1.26e+04
x38	1.173e+04	5042.938	2.326	0.020	1846.790	2.16e+04
x39	-1439.6685	6367.366	-0.226	0.821	-1.39e+04	1.1e+04
x40	2.018e+04	5731.869	3.521	0.000	8948.092	3.14e+04
x41	20.7011	5.315	3.895	0.000	10.283	31.120
x42	0.5203	6.149	0.085	0.933	-11.533	12.574
x43	497.0107	166.247	2.990	0.003	171.139	822.882
x44	-29.4154	8.471	-3.472	0.001	-46.021	-12.810
x45	0.2802	0.250	1.121	0.262	-0.210	0.770
x46	254.8610	4790.816	0.053	0.958	-9135.948	9645.670
x47	374.5480	436.920	0.857	0.391	-481.889	1230.985
x48	-57.9071	16.045	-35.609	0.000	-89.357	-26.457
x49	-190.0649	52.656	-3.610	0.000	-293.280	-86.850
x50	1.8153	6.928	0.262	0.793	-11.764	15.394
x51	174.6081	18.263	9.561	0.000	138.809	210.407

Omnibus:	4244.626	Durbin-Watson:	2.011
Prob(Omnibus):	0.000	Jarque-Bera (JB):	81692.861
Skew:	1.256	Prob(JB):	0.00
Kurtosis:	15.701	Cond. No.	1.07e+16

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.07e+16. This might indicate that there are strong multicollinearity or other numerical problems.

[]: